

Establishing trust in automated reasoning

Konrad Hinsen

konrad.hinsen@cnrs.fr

Centre de Biophysique Moléculaire (UPR4301 CNRS)
Rue Charles Sadron, 45071 Orléans Cédex 2, France

Synchrotron SOLEIL, Division Expériences
B.P. 48, 91192 Gif sur Yvette, France

Abstract

Since its beginnings in the 1940s, automated reasoning by computers has become a tool of ever growing importance in scientific research. So far, the rules underlying automated reasoning have mainly been formulated by humans, in the form of program source code. Rules derived from large amounts of data, via machine learning techniques, are a complementary approach currently under intense development. The question of why we should trust these systems, and the results obtained with their help, has been discussed by early practitioners of computational science, but was later forgotten. The present work focuses on independent reviewing, an important source of trust in science, and identifies the characteristics of automated reasoning systems that affect their reviewability. It also discusses possible steps towards increasing reviewability and trustworthiness via a combination of technical and social measures.

1 Introduction

Like all social processes, scientific research builds on trust. In order to increase humanity's knowledge and understanding, scientists need to trust their colleagues, their institutions, their tools, and the scientific record. Moreover, science plays an increasingly important role in industry and public policy. Decision makers in these spheres must therefore be able to judge

which of the scientific findings that matter for them are actually trustworthy.

In addition to the trust-forming mechanisms present in social relationships, the scientific method is built in particular on *transparency* and *independent critical inspection*, which serve to remove mistakes and biases in individual contributions as they enter the scientific record. Ever since the beginnings of organized science in the 17th century, with the appearance of learned societies and the first scientific journals, researchers are expected to put all facts supporting their conclusions on the table, and allow their peers to inspect them for accuracy, pertinence, completeness, and bias. Since the 1950s, following a sharp increase in the number of researchers and the availability of photocopiers, critical inspection in the form of *peer review* has become an integral part of the publication process [Spier 2002]. It is still widely regarded as a key criterion for trustworthy results.

Over the last two decades, an unexpectedly large number of peer-reviewed findings across many scientific disciplines have been found to be irreproducible upon closer inspection [Baker 2016; Stodden et al. 2018; Colliard et al. 2022; Samuel and Mietchen 2024]. This so-called “reproducibility crisis” has shown that our practices for performing, publishing, reviewing, and interpreting scientific studies are no longer adequate in today’s scientific research landscape, whose social, technological, and economic contexts have changed dramatically. Updating these processes is a major aspect of the nascent Open Science movement.

The topic of this article is a particularly important recent change in research practices: the increasing use of automated reasoning. Computers and software have led to the development of completely new techniques for scientific investigation, and permitted existing ones to be applied at larger scales and by a much larger number of researchers. In the quantitative sciences, almost all of today’s research critically relies on computational techniques, even when they are not the primary tool for investigation [Hettrick et al. 2014; Nangia and Katz 2017]. Simulation, data analysis, and statistical inference have found their place in almost every researcher’s toolbox. Machine learning techniques, currently under intense development, are likely to become equally ubiquitous in the near future.

From the point of view of transparency and critical inspection, these new tools are highly problematic. In philosophy of science, this lack of transparency is referred to as *epistemic opacity* [Humphreys 2009]. Among practitioners of computational science, it was a topic of debate in the 1980s, when affordable desktop workstations first made computation accessible to a larger number of scientists. It has been summarized by Turkle [2009] as “the tension between doing and doubting”, i.e. the tension between enthusiasm for the new possibilities and the doubts about the reliability of the very new and unproven techniques. Turkle also describes the subsequent waning of doubt:

“Familiarity with the behavior of virtual objects can grow into something akin to trusting them, a new kind of witnessing.” Today, a typical research paper reports the use of software uncritically, much like the use of well-understood scientific instruments. There are rarely any signs of reflection about the reliability of the software or about its adequacy for the task being performed.

Automation bias [Parasuraman and Riley 1997], i.e. the propensity to consider automated processes more reliable than human labor, is one possible cause. Resignation to the impossibility of constructive doubt is another one. The results produced by automated reasoning are often neither obviously correct nor obviously wrong. In the absence of process transparency, this leaves only uncritical acceptance or uncritical rejection as possible reactions. Today’s default is uncritical acceptance.

The reproducibility crisis could have been a wakeup call, but the contribution of automated reasoning to this crisis has not been widely recognized. Several examples cited in this context involve faulty software [Merali 2010], e.g. the retraction of papers in structural biology due to a software bug [Miller 2006], or the mistakes discovered in a high-impact study on the relation between debt and economic growth [Herndon et al. 2014]. These documented cases are probably only the tip of the iceberg, given that detecting a mistake and tracing it back to a software issue is an arduous task when all elements have been published, and an impossible one otherwise. Pessimistic but not unrealistic estimates suggest that most computational results in science are to some degree wrong because of software defects [Soergel 2015; Thimbleby 2023].

However, most high-profile cases of observed non-reproducibility involve empirical research, e.g. in medicine or psychology, in which automated reasoning is not seen as a major aspect because it intervenes only in the routine application of statistical software for data analysis. Nevertheless, there are two aspects of this use of software that deserve critical reflection. First, incorrect use of software is as much as source of mistakes as defective software. Given the bad state of user interface design and documentation for much scientific software, incorrect use should be expected to happen. Second, inappropriate use is an even more subtle source of errors. The mere existence of statistical software has made statistical methods accessible to a much wider public. Techniques that used to require collaboration with trained statisticians are nowadays available as black-box tools to researchers who may not even know what they would need to learn in order to use these techniques correctly.

Given that it is impossible in practice to follow a computation step by step, meaning processor instruction by processor instruction, in order to verify its correctness, on what basis can scientists justify their belief in its results? Durán and Formanek [Durán and Formanek 2018] propose *computational*

reliabilism as a general framework. It stipulates that a result can be trusted if the algorithm or piece of software used to obtain it can be considered reliable. Reliability is evaluated using various reliability indicators, which can be based on characteristics of the source code but also on the observed behavior of the algorithm and on the context of its actual use in scientific practice [Durán 2025]. The choice of suitable reliability indicators, as well as the weight that should be attributed to each of them, remains an open issue and should be expected to be context-dependent. From a pragmatic rather than epistemologic point of view, reliabilism is very close to how scientist actually reason about their software, once they start reflecting on it.

From the perspective of computational reliabilism, critical inspection has the goal of evaluating appropriate reliability indicators for each piece of software. It can involve reading the source code and its documentation, computing software quality metrics, testing programs on well-known test cases, performing usability studies, etc. Given today’s complex software stacks, an exhaustive critical inspection is impossible for the individual scientist wishing to deploy it, for lack of both time and competence. The division of labor that makes the development of complex software stacks possible requires a corresponding division of labor in critical inspection. Individual components as well as larger assemblies must be reviewed by independent experts, both in order to provide feedback to software authors and in order to allow software users to evaluate the reliability of their tools based on expert judgment. An added complication is that reliability is contextual, because the software’s behavior can be appropriate for one use case but unacceptable for another one. A review must therefore either be exhaustive or clearly state the limits of its focus.

Open Science has made a first step towards dealing with automated reasoning, in insisting on the necessity to publish scientific software, and ideally making the full development process transparent by the adoption of Open Source practices. While this level of transparency greatly facilitates critical inspection, it is not sufficient. Someone must actually look at this software with a critical eye, looking for mistakes and for tacit assumptions that users need to know about. This is not happening today, and we do not even have established processes for performing such reviews [Hinsen 2025]. Moreover, as I will explain later, much of today’s scientific software is written in a way that makes independent critical inspection particularly challenging if not impossible. If we want scientific software to become trustworthy, we therefore have to develop reviewing practices in parallel with software architectures that make reviewing actually feasible in practice. And where reviewing is not possible, we must acknowledge the fragile nature of automated reasoning processes and make sure that everyone looking at their results is aware of their uncertain reliability.

As for all research tools, it is not only the software itself that requires critical inspection, but also the way the software is used in a specific research project. Improper use of software, or inappropriateness of the methods implemented by the software, can lead to mistakes as well. However, it is much more difficult to detect with today’s minimal reporting practices. Moreover, the distinction between a defect and inappropriate use is not as obvious as it may seem. A clear distinction would require a well-defined interface between software and users, much like a written contract. If the software’s behavior deviates from this contract, it’s a defect. If the user’s needs deviate from the contract, it’s inappropriate use. But such detailed contracts, called *specifications* in the context of software, rarely exist. The high cost of writing, verifying, and maintaining specifications limits their use to particularly critical applications. This means that reviewing the use of scientific software requires particular attention to potential mismatches between the software’s behavior and its users’ expectations, in particular concerning edge cases and tacit assumptions made by the software developers. They are necessarily expressed somewhere in the software’s source code, but users are often not aware of them.

The scientific requirement of *independent* reviewing is related to another aspect of automated reasoning that I will address, in particular in my proposals for improving our current practices: the preservation of epistemic diversity. As Leonelli has pointed out [Leonelli 2022], the Open Science movement has so far largely neglected this aspect. Epistemic diversity is about different perspectives and research methodologies coexisting, enriching and critiquing each other. Automation, be it in industry or in research, tends to reduce diversity by favoring standardization as an enabler of economies of scale. In the Open Science movement, this tendency is implicit in the quest for reusability, one of the four FAIR principles [Wilkinson et al. 2016; Barker et al. 2022]. Reusing someone else’s code or data requires adopting the authors’ methodologies, and to some degree their general perspective on the phenomenon under study. In the extreme case of a single software package being used by everyone in a research community, there is nobody left who could provide critical feedback.

This article has two main parts. In the first part (section 2), I look at the factors that make automated reasoning more or less reviewable. It is a critical examination of the state of the art in scientific software and its application, which should help scientists to get a better grasp of how reliable automated reasoning can be expected to be. In the second part (section 3), I consider how the reviewability of automated reasoning can be improved, both through better reviewing processes and by restructuring software for better reviewability.

2 Reviewability of automated reasoning systems

Automated reasoning can play different roles in scientific research, with different reliability requirements.¹ The numerical preprocessing of observational data before scientific analysis, nowadays often integrated into scientific instruments, is an example where high reliability is required, because its outputs are used without any further verification. On the other hand, protein structure prediction by AlphaFold [Jumper et al. 2021] is known to be unreliable, but it is nevertheless very useful if coupled with experimental validation of its predictions [Nielsen 2023]. Traditional computer simulation is often used similarly in biology as a hypothesis generator whose outputs are subject to subsequent validation, whereas in engineering, simulations of mechanical systems are routinely performed to support critical decisions, thus requiring high reliability.

What these examples illustrate is that tools, processes, and results in science do not necessarily have to be perfectly reliable. Higher-level validation processes act much like error correction protocols in engineering. The coherence of multiple approaches to a question, coming from different perspectives, is another higher-level source of reliability, indicating robustness. This again illustrates the importance of epistemic diversity that I have mentioned in the introduction. What matters, however, is a clear understanding of the reliability of individual scientific contributions, which in turn requires a clear understanding of the reliability of the tools and processes on which those contributions are based.

In this section, I discuss five characteristics (summarized in Fig. 1) of automated reasoning systems that influence how their reliability can be assessed by independent critical inspection, which in the following I will call *review* for brevity. This use of *review*, inspired by the tradition of scientific peer review, should not be confused with the software engineering technique of *code review*, which is a quality control step performed *internally* by a development team. The term should not be read either as a direct transposition of the pre-publication peer review of papers by scientific journals, whose main output is an accept-or-reject judgement. Reviewing can take various forms depending on context. It can be performed by the researchers who wish to adopt someone else’s software tool, or by designated experts whose reports are published.

Also for brevity, I will use the term *software* instead of “automated reasoning system”, extending its usual meaning to include trained neural networks and other models obtained via machine learning techniques. In science, software and machine learning differ in degree rather than in kind, as already long before machine learning, many computational models contained a large

¹This is of course true for software in general, see e.g. the discussion in [Shaw 2022:22].

number of parameters fitted to large datasets.

Wide spectrum	—————	Situated
Stable	—————	Evolving
Transparent	—————	Opaque
Convivial	—————	Open
	—————	Proprietary
Few dependencies	—————	Many dependencies

Figure 1: The five dimensions of scientific software that influence its reviewability.

Recently, Hasselbring *et al.* [2025] have published a multidimensional categorization of research software that doesn’t address reviewability, but has categories that overlap partially with the ones I define in the following. Readers unfamiliar with the diversity of research software would profit from reading their analysis before continuing here.

Fig. 2 shows the layers of a typical software stack for a research project, which I will refer to in the following. The top layer is software written specifically for a research project. This layer usually contains small items: scripts, notebooks, or workflows. They use software components from the layer below, which I have labelled “domain-specific tools”. This domain-specific layer contains software that implements the computational models and methods used by a research community. Many ongoing discussions of scientific software, in particular concerning its sustainability [Hettrick 2016], concentrate on this layer, but are not always explicit about this focus.

The two levels further down contain infrastructure software, meaning software that the domain-specific layer builds on but which computational scientists are often only vaguely aware of. They may for example know that their software is written in the C language, without knowing which precise version of which C compiler was used to compile it. I have divided infrastructure into two layers: a general one that computational science shares with other application domains of computing, and a scientific layer for software that is written explicitly to support science and engineering.

2.1 Wide-spectrum vs. situated software

Wide-spectrum software provides fundamental computing functionality to a large number of users. In order to serve a large user base, it addresses a wide range of application scenarios, each of which requiring only a part of the software’s functionality. Word processors are a well-known example: a

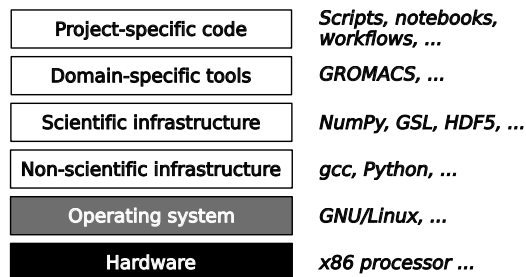


Figure 2: A typical software stack as used in a research project

package like [LibreOffice](#) can be used to write a simple letter, but also a complex book. LibreOffice has huge menus filled with various functions, of which most users only know the handful that matters to them. General-purpose large language models are another example of wide-spectrum software.

Situated software (a term introduced by Shirky [2004]) is software written for a specific use case or a narrow user group. It addresses a specific need very well, but is not transferable to other application scenarios. Spreadsheets are usually situated, as are games, and many shell scripts.

A useful numerical proxy for estimating a software package’s location on this scale is the ratio of the number of users to the number of developers, although there are exceptions. Games, for example, are situated software with few developers but many users.

In scientific computing, the wide-spectrum end of the scale is well illustrated by mathematical libraries such as [BLAS](#) or visualization libraries such as [matplotlib](#), which provide a large collection of functions from which application developers pick what they need. They are found mostly in the infrastructure layer of Fig. 2, but also in the domain-specific layer. At the situated end, we have mainly the project-specific code at the top of Fig. 2: code snippets and scripts that generate the plots shown in a paper, or computational notebooks and workflows. Most of the domain-specific layer lies in between these extremes; see section 2.7.3 below for a case study.

Reviewing wide-spectrum software represents a major effort, because of its size and functional diversity. Moreover, since wide-spectrum software projects tend to be long-lived, with the software adapting to new use cases and new computing platforms, its critical examination must be an ongoing process as well. On the other hand, this effort can be considered a good investment, because of the large user base such software has.

Situated software is usually smaller and simpler, which makes it easier to understand and thus to review. However, its evaluation can only be done in the specific context for which the software was written. This suggests

integrating it into the existing scientific peer reviewing process, along with papers and other artifacts that result from a research project.

It is the intermediate forms of software that are most difficult to review. Domain tools and libraries are too large and complex to be evaluated in a single session by a single person, as is expected in today’s journal peer review process. However, they don’t have a large enough user base to justify costly external audits, except potentially in contexts such as high-performance computing where the importance of the application and the high cost of the invested resources also justify more careful verification processes.

2.2 Stable vs. evolving software

Stable software is developed and maintained with the goal of providing a reliable tool. Signs of stability in software are its age, a clear definition of its purpose, ideally in the form of a specification, respect of standards, respect of software engineering practices, detailed documentation, and the absence of compatibility-breaking changes. The Linux kernel and the text editor Emacs are examples of very stable software.

Stable does not mean static. It is common for stable software to grow by addition of new functionality. It is also common for stable software to be ported to new computational environments. Stability is evaluated from the user’s point of view, not by counting changes to the source code. Both the documented behavior of the software and the tacit knowledge that users acquire over the years should be stable.

Evolving software is either software in early stages of development, for which design and implementation decisions are still revised from time to time, or software whose goals change over time. One important category of evolving software is prototype software, i.e. software created to test new ideas, be they technical (software architecture etc.) or related to the application domain. For software projects that adopt [semantic versioning](#), the major version number 0 indicates prototype status.

Another important category is software that is continuously adapted to changing requirements. This has become common over the last two decades, in particular within Open Source communities. As an illustration of this rising attitude, consider programming languages: all languages with a written standard and multiple implementations are more than 20 years old (but the standards are still updated from time to time), whereas programming languages published more recently have a single authoritative implementation that acts as an evolving *de facto* specification of their behavior. This does not prevent them from being stable, by consensus of the developer community, as e.g. for the Clojure programming language [Pote 2025], but stability cannot be taken for granted, even for a 30 year old language such as Python.

Unfortunately, developers usually make no clearly visible statement about aiming for stability or for continuous adaptation.

Research software can be anything from very stable to rapidly evolving. Many domain-specific software packages (see Fig. 2) contain both stable and evolving parts. For their developer communities, which are typically small, keeping stable and evolving parts separate is neither easy nor particularly important, because the potential users for both parts are the same. Infrastructure ought to be stable in order to ensure a solid foundation for a software stack. This is generally true for infrastructure software having its roots in the 20th century. Many more recent software packages are *de facto* used as infrastructure in spite of being evolving. See section 2.7.1 below for a case study.

For reviewing, evolving software is mainly an economic issue. Reviewing takes time, and the reviews must evolve at the same pace as the software itself. But from the wider perspective of reliability and trust, evolving software raises the more fundamental issue of what reliability even means for a piece of software whose intended behavior is in flux. Building trust of any kind takes time. It requires repeated probing with dominantly positive returns. People will never see software as reliable if it evolves faster than they can build trust in it. This implies that reliability is relational, i.e. not an absolute characteristic of a piece of software but a judgement made from a specific perspective. The same software package can be seen as reliable by a power user, but as unstable by an occasional user.

Finally, evolving software requires more attention from its users, who must ensure that their knowledge about the software is up to date. The helpful hint found on your colleague's Web site may no longer be appropriate. Your own well-documented use case from a few years ago may no longer be a good template to start a new project. Evolving software thus comes with a higher cognitive load and an increased risk of inappropriate use.

2.3 Transparent vs. opaque software

Transparent software is software whose behavior is comprehensible and verifiable by its users. Transparency differs from openness in that it is about the behavior of a program, whereas openness is about its source code. In a word processor, or a graphics editor, many user actions, such as pasting in a piece of text or an image, produce an immediately visible result that the user can easily predict and verify. In contrast, opaque software performs complex operations and produces output whose correctness cannot be assessed with reasonable effort. Large language models are an extreme example: their output is impossible to predict and verify. Note that this notion of opacity is not identical but similar to the epistemic opacity I

discussed in the introduction.

Strictly speaking, transparency is not a characteristic of a piece of software, but of a computational task. A single piece of software may contain both transparent and opaque functionality. In the word processor example, inserting a character is highly transparent, whereas changing the page layout is more opaque, creating the possibility of subtle bugs whose impact is not readily observable. I use the term “opaque software” as a shorthand for software implementing at least one major opaque operation.

In scientific software, transparency often comes from a simple specification for which multiple implementations are available. Examples are solvers for partial differential equations in fluid dynamics, quantum chemistry computations, or common numerical libraries for linear algebra or fast Fourier transforms. The existence of multiple implementations of precisely the same task makes it possible to verify results from one implementation by comparing to the output of other implementations. Moreover, the simplicity of the specification makes it possible to perform tests on inputs for which the exact solution is known by different means.

Most scientific software is closer to the opaque end of the spectrum. Even highly interactive software, for example in data analysis, performs non-obvious computations, yielding output that an experienced user can perhaps judge for plausibility, but not for correctness. Statistical inference processes are a good example. As a rough guideline, the more scientific models or observational data have been integrated into a piece of software, becoming a tacit part of its specification, the more opaque the behavior of the software is to its users.

It is much easier to evaluate the reliability of transparent software, because its expected output can be obtained independently. Reviewing transparent software is therefore easier, to the point that it can often be done by the user rather than by an expert. Moreover, when users can understand and judge the results produced by a piece of software, even a very weak reliability indicator such as popularity becomes useful.²

The more opaque a computation is, the more important its documentation becomes. Reliability can only be judged for the performance of a well-defined task. If users cannot know what the software promises to do exactly, nor what its limitations are, then they cannot perform any task reliably using this software. This is currently a much discussed issue with machine learning models [e.g. Gunning et al. 2019; Longo et al. 2024], but it is not sufficiently

²A famous dictum in software engineering, often referred to as “[Linus’ law](#)”, states that “given enough eyeballs, all bugs are shallow”. However, this can only work if the many eyeballs are sufficiently trained to spot problems, meaning that users of opaque software don’t qualify.

recognized that traditional computer software can be just as opaque from a user’s point of view, if source code is the only available documentation of its behavior.

2.4 Convivial vs. proprietary software

Convivial software [Kell 2020] is named in reference to Ivan Illich’s book “Tools for conviviality” [Illich 1973]. It draws on the root meaning of the latin *convivium*, which is “living together”. Convivial technology as described by Illich is shaped by the community of its users according to their evolving needs, rather than being controlled by specialists outside of the community. Convivial software aims at augmenting its users’ agency over their computations by actively encouraging them to adapt the software to their needs. [Malleable software](#) is a similar concept, as is re-editable software, a term introduced by [Donald Knuth in an interview](#) in opposition to reusable, i.e. off-the-shelf, software [Hinsen 2018a].

In contrast, proprietary software offers users fixed functionality and therefore limited agency. At first sight this looks like users should always prefer convivial software, but agency comes at a price: users have to invest more effort and assume responsibility for the final software assembly that they have made themselves. In fact, conviviality is all about diminishing the gap between software developers and software users. Just like most people prefer to choose from a range of industrially-made refrigerators, rather than modify a generic design precisely to their needs, most computer users are happy to use ready-made e-mail software rather than assembling their own from items in a toolkit.

It is important to understand that conviviality is not a characteristic of software as a digital artifact (e.g. the contents of a repository on GitHub), but of a software ecosystem as a socio-technical system. A well-known convivial ecosystem is the collection of Unix command line tools, which users can freely combine in shell scripts and complement by adding their own tools. It doesn’t make much sense to call individual command line tools such as `ls` or `grep` convivial, but it does make sense to say that they are designed to be part of a convivial ecosystem. A necessary but not sufficient condition for fitting into a convivial ecosystem is transparency (see [section 2.3](#) and [sections 2.7.1](#) and [2.7.3](#) for relevant case studies).

In the academic literature on software engineering, convivial software is discussed with the focus on its users becoming developers, most commonly referred to as *end user programmers* [Nardi 1993; Ko et al. 2011]. Shaw recently proposed the less pejorative term *vernacular developers*, which she defines as “people who are not professionally trained as programmers [and who] are creating and tailoring software as a means to achieve goals of their

own” [Shaw 2022]. The subfield of end user software engineering aims at providing vernacular developers with methods and tools to improve the quality of their software, recognizing that the methods and tools designed for software professionals are usually not adapted to their needs.

The risk of non-convivial technology, which Illich [1973] discusses in detail, is that widespread adoption makes society as a whole dependent on a small number of people and organizations who control the technology. This is exactly what has happened with computing technology for non-professional use, such as personal computers and smartphones. You may not want to let corporations spy on you via your smartphone, but the wide adoption of these devices means that you are excluded from more and more areas of social life if you decide not to use one. Many research communities have fallen into this trap as well, by adopting proprietary infrastructure, such as [MATLAB](#), or Open Source infrastructure over which they have no control, such as Python, as a foundation for their computational tools and models. A convivial infrastructure would have preserved these researchers’ autonomy to shape the software according to the needs of their research, at the cost of much stronger participation in its development and maintenance.

In between convivial and proprietary software, we have Free, Libre, and Open Source software (FLOSS). Historically, the Free Software movement was born in the 1980s in academia [Gonzalez-Barahona 2021]. Much software was distributed with its source code, even in commercial settings. Most computer users in this environment needed to do some programming of their own in order to get anything done. In other words, software was much more convivial than what we have today. It is the arrival of proprietary software in their lives, exemplified by the frequently cited proprietary printer driver at MIT [2002], that pushed them towards formalizing the concept of Free Software in terms of copyright and licensing.

With the enormous complexification of software over the following decades, a license is no longer sufficient to keep software convivial in practice. The right to adapt software to your needs is of limited value if the effort to do so is prohibitive. Software complexity has led to a creeping loss of user agency, to the point that even building and installing Open Source software from its source code is often no longer accessible to non-experts because of the complexity of the build tools. An experience report on building the popular machine learning library [PyTorch](#) from source code nicely illustrates this point [Courtès 2021]. The authors of PyTorch explain how you can install the precompiled versions they provide (via Python’s `pip` utility), but they do not explain how to compile the code on your own, as you would have to do to fix bugs or tweak some functionality. This makes it practically impossible for most PyTorch users to adapt the code to their needs. The work by Courtès [2021] amounts to reverse-engineering the build process and publishing it as

a build recipe for [Guix](#), a package manager designed for reproducible builds from source code.

Conviviality has become a marginal subject in the FLOSS movement, with the Free Software subcommunity pretending that it remains ensured by copyleft licenses and much of the Open Source subcommunity not caring about it. It survives mainly in communities whose technology has its roots in the 1970s and 1980s, when ideas such as Alan Kay’s Dynabook [Kay 1972], a device for which even children could create software for their own use, were active research topics. Examples are programming systems inheriting from Smalltalk (e.g. [Squeak](#), [Pharo](#), and [Cuis](#)), or the programmable text editor [GNU Emacs](#). These systems encourage modification via built-in tools for discovering functionality and for locating the parts of the code that implement it.

In scientific computing, there is a lot of diversity on this scale. Fully proprietary software is common, but also variants that do allow users to look at the source code, but don’t allow them to compile it, or don’t allow the publication of reviews. In computational chemistry, the widely used Gaussian software is an example for such legal constraints [Hocquet and Wieber 2017]. FLOSS has been rapidly gaining popularity, and receives strong support from the Open Science movement (see e.g. the Second French National Plan for Open Science [French Ministry for Higher Education, Research, and Innovation 2021] that has a whole chapter dedicated to software). Somewhat surprisingly, the move beyond FLOSS to convivial software is hardly ever envisaged, in spite of it being aligned with the pre-digital practices of scientific research: the main intellectual artifacts of science, i.e. theories and models, have always been convivial.

Concerning reviewing, the convivial-to-open part of the scale is similar to the situated-to-wide-spectrum scale: convivial software is easier to understand and therefore easier to review, but each specific adaptation of convivial software requires its own review, whereas open but not convivial software makes reviewing a better investment of effort. Proprietary software is harder to review, because only its observed behavior and its documentation are available for critical inspection, but not its source code.

2.5 Size of the minimal supporting environment

Each piece of software requires a computational environment, consisting of a computer and of other pieces of software. In terms of the software stack illustrated by Fig. 2, the computational environment for a piece of software consists is situated in the layers below its own. The importance of computational environments is not sufficiently appreciated by most researchers today, which is the main cause of widespread computational irreproducibility

[Stodden and Miguez 2013; Hinsén 2020; Wang et al. 2021].

It is the computational environment that defines what a piece of source code actually does. The meaning of a Python script is defined by the Python interpreter. The Python interpreter is itself a piece of software written in the C language, and therefore the meaning of its source code is defined by the C compiler, and by the processor which ultimately executes the binary code produced by the C compiler. As an illustration for the importance of the computational environment, it is an easy exercise to write a Python script that produces different results when run with version 2 or version 3 of the Python interpreter, exploiting the different semantics of integer division between the two versions.

In addition to this semantic importance of computational environments, reviewability implies a pragmatic one: reviewers of software or its results need access to an adequate hardware and software environment in order to perform their review. Scientific computing mostly relies on commodity hardware today, with two important exceptions: supercomputers and Graphical Processing Units (GPUs). Supercomputers are rare and expensive, and thus not easily accessible to a reviewer. GPUs are evolving rapidly, making it challenging to get access to an identical configuration for reviewing. Supercomputers often include GPUs, combining both problems. Resource access issues are manageable for wide-spectrum software if they are deemed sufficiently important to warrant the cost of performing audits on non-standard hardware.

Software environments have been recognized as highly relevant for automated reasoning in science and beyond, and are the subject of active research [e.g. Malka et al. 2024; Bilke et al. 2025]. They play a key role in computational reproducibility, but also for privacy and security, which are the prime motivations for the [Reproducible Builds](#) movement ³. The issues of managing software environments are now well understood, and two software management systems ([Nix](#) and [Guix](#)) implement a comprehensive solution. However, they are so far used only by a small minority of researchers. In addition to ease of use issues, which could be overcome with time and investments, a major obstacle is that such management systems must control the complete software stack, which excludes the use of popular proprietary platforms such as Windows ⁴ or macOS ⁵.

Assuming that the proper management of scientific software environments will be achieved not only in theory, but also in practice, it is the size of

³Ken Thompson’s Turing Award Lecture “Reflections on Trusting Trust” [Thompson 1984] is an early and very readable discussion of the security implications of computational environments.

⁴Windows is a trademark of the Microsoft group of companies.

⁵macOS is a trademark of Apple Inc., registered in the U.S. and other countries and regions.

this environment that remains a major characteristic for reviewability. The components of the computational environment required by a piece of software are called its *dependencies* in software engineering. This term expresses their importance very well: every single quality expected from a software system is influenced by the corresponding quality of the components that enter in its construction. For example, no software can be more stable than its dependencies, because of the risk of software collapse [Hinsen 2019]. Reviewing software therefore requires an examination of its dependencies as well. This can become an obstacle for software that has hundreds or even thousands of dependencies. Imprecise dependency references (e.g. “version 2 or later of package X”) are a problem as well, because they require reviewers to emit a judgement on a potentially large number of versions of a package, and to extend their trust in them to an indefinite future.

2.6 Analogies in experimental and theoretical science

For developing a better understanding of the reviewability characteristics described above, it is helpful to consider analogies from the better understood experimental and theoretical techniques in scientific research. In particular, it is helpful to examine where such analogies fail due to the particularities of software.

Experimental setups are situated. They are designed and constructed for a specific experiment, described in a paper’s methods section, and reviewed as part of the paper review. Most of the components used in an experimental setup are mature industrial products, ranging from commodities (electricity cables, test tubes, etc.) to complex and specialized instruments, such as microscopes and NMR spectrometers. Non-industrial components are occasionally made for special needs, but this is discouraged by their high manufacturing cost. The use of prototype components is exceptional, and usually has the explicit purpose of testing the prototype. Some components are very transparent (e.g. electricity cables), others are very opaque (e.g. NMR spectrometers). The equivalent of the computational environment is the physical environment of the experimental setup. Its impact on the observations tends to be well understood in the physical sciences, but less so in the life sciences, where it is a common source of reproducibility issues (e.g. [Kortzfleisch et al. 2022] or [Georgiou et al. 2022]).

The main difference to software is thus the much lower prevalence of prototype components. A more subtle difference between instruments and software is that the former are carefully designed to be robust under perturbations, whereas computation is chaotic [Hinsen 2016]. A microscope with a small defect may show a distorted image, which an experienced microscopist will recognize when evaluating the microscope. Software with a small defect, on the other hand, can introduce unpredictable errors in both kind and

magnitude, which neither a domain expert nor a professional programmer or computer scientist can diagnose easily if the software is opaque. This is a consequence of software being crafted out of a substrate (Turing-complete languages) that is much less constraining than the physical substrates from which instruments are made. The increasing integration of computers and software into scientific instruments may lead to experimental setups becoming less robust as well in the future.

Analogies with pre-digital scientific models and theories are instructive as well. Wide-spectrum theories exist in the form of abstract reasoning frameworks, in particular mathematics. The analogue of situated software are concrete models for specific observational contexts. In between, we have general theoretical frameworks, such as evolutionary theory or quantum mechanics, and models that intentionally capture only the salient features of a system under study, pursuing understanding rather than precise prediction. Examples for the latter are the Ising model in physics or the Lotka-Volterra equations in ecology.

Abstract frameworks and general theories are the product of a long knowledge consolidation process, in which individual contributions have been reviewed, verified on countless applications, reformulated from several perspectives, and integrated into a coherent whole. This process ensures stability, transparency, and conviviality in a way that has so far no equivalent for software.

Opacity is an issue for theories and models as well: they can be so complex and specialized that only a handful of experts understand them. It also happens that people apply such theories and models inappropriately, for lack of sufficient understanding. However, automation via computers has amplified the possibility to deploy opaque sets of rules so much that it makes a qualitative difference: scientists can nowadays use software whose precise function they could not understand even if they dedicated the rest of their career to it.

The computational environment for theories and models is the people who work with them. Their habits, tacit assumptions, and metaphysical beliefs play a similar role to hardware and software dependencies in computation, and they are indeed also a common cause of mistakes and misunderstandings.

2.7 Case studies

In order to illustrate the five dimensions described above, I will present a few case studies on software that I know well, from having used it in my own research work.

2.7.1 NumPy

NumPy [Harris et al. 2020] is the foundational library of the Scientific Python ecosystem. It adds an efficient array data structure to the Python language, which serves both for efficient manipulation of scientific data from Python and for interfacing Python with lower-level languages such as C or Fortran. NumPy is an interesting case study because of its status as widely used infrastructure. It is also an example I know exceptionally well because I was involved in its early development stages [Dubois et al. 1996].

Scope: NumPy is used in all disciplines of quantitative science, for data of very diverse types, including geometric objects, pixel images, and neural networks. It is therefore clearly a wide-spectrum software package.

Openness: NumPy is Open Source software. It is the foundation of an ecosystem that is not convivial, being much too complex for its users to understand or modify. However, its transparency (see below) would make it a candidate for a foundation of a convivial research software ecosystem, if it could be made more stable (see below).

Stability: The stability status of NumPy is what makes it an interesting case study. NumPy grew out of its predecessor *Numerical Python*, whose development started in the 1990s [Dubois et al. 1996]. It was very stable in its first decade, with the library and its dependees growing but maintaining backward compatibility. In the 2000s, the popularity of the scientific Python ecosystem grew rapidly. In 2006, NumPy was published as a major redesign that unified the original Numerical Python with its similar but not fully compatible offspring *numarray*. At this occasion, numerous changes were made to the API to ease the transition from Matlab to Python, corresponding to the needs of a significant part of the freshly recruited community members. The old API was maintained under a different name in order to make migration from Numerical Python to NumPy straightforward and risk-free. The intention was to let the two API versions live side by side indefinitely. But another few years later, after more growth, the old API was declared obsolete and then removed, by younger maintainers who had no memories of it anyway. This broke some of the oldest scientific software applications that had been written in Python. However, their developers and users were no longer representative for the community, whose culture had shifted to valuing innovation over stability, in alignment with the wider technology culture [Hinsen 2024b]. Since then, NumPy has regularly introduced breaking changes in order to make way for what the developers of the moment considered improvements, culminating in a recent major-version update to NumPy 2 that comes with an extensive [migration guide](#).

Transparency: Basic array operations are very transparent. However, there are many of them, and some have non-trivial edge cases. Verifying the

correct behavior of all operations is not difficult but laborious. An extensive collection of unit tests with good coverage takes care of that.

Dependencies: Identifying the dependencies of NumPy is not an easy exercise. The installation instructions claim that “the only prerequisite for installing NumPy is Python itself.” However, the detailed instructions for compiling the source code add compilers for C and C++, the libraries BLAS and LAPACK, and pkg-config. All of these packages have their own dependencies as well. The complete dependency graph for NumPy 2.3.1 as recorded by the [package manager Guix](#) in commit `da04b3bb42d76e7af9d5eb344cfde2350e9bb3c1` contains 239 items (including NumPy itself) ⁶. Most of these packages are well-known and widely used stable elements of common computational infrastructure, the notable exception being the Python language, which has been evolving in a way similar to NumPy itself.

Its evolving nature makes NumPy difficult to review, but on the other hand its transparency makes reviewing less important, because software engineering practices already provide sufficient reliability indicators. However, the evolving nature of NumPy (and its main dependency, Python) is a significant obstacle to reviewing software that builds on it. That is an enormous amount of software across many disciplines. NumPy is *de facto* part of today’s scientific computing infrastructure. Infrastructure ought to be stable, but NumPy and Python are not.

A major contribution to this unsatisfactory situation is the lack of support for computational infrastructure in general. Infrastructure maintenance requires institutions that oversee and fund it from a long-term perspective, and with a governance model in which all types of users are represented. Such institutions do not exist today. As a consequence, infrastructure such as NumPy is developed according to shifting criteria that represent the interests of whoever happens to be funding the current generation of developers.

2.7.2 GNU Scientific Library (GSL)

The GNU Scientific Library [Galassi 2009; 2025] implements widely used algorithms for numerical mathematics in the C language, with interfaces for many other languages. It has been under continuous development since 1996. GSL is an interesting case study in comparison with NumPy, because it is also infrastructure software but very stable and more convivial.

Scope: GSL is a wide-spectrum library, containing functionality of interest to many domains of quantitative science.

⁶To reproduce this number, run `guix time-machine --commit=da04b3bb42d76e7af9d5eb344cfde2350e9bb3c1 -- graph python-numpy | grep label | wc -l`

Openness: GSL is Free Software (GPL license), and thus open. It is even halfway towards convivial. Being a collection of independent or loosely coupled modules, it makes modifying a single module quite accessible. When used as a component of a software stack, GSL is typically situated in a lower layer, making it just as difficult to replace with a modified version as any other lower-layer component. But it is feasible to extract a GSL module and add a modified copy of it, under a different name, to the software stack again.

Stability: A look at the [summary of changes](#) shows that GSL is very stable. Changes fix bugs, add new functionality, and add support for more platforms.

Transparency: Being a collection of implementations of well-known algorithms of moderate complexity, GSL is quite transparent.

Dependencies: The only dependency for using GSL is a C compiler. The [BLAS](#) implementation included with GSL can be replaced by a more efficient one, but this is optional. Modifying and then rebuilding GSL, including its documentation, requires five few well-known and stable build tools: [autoconf](#), [automake](#), [libtool](#), [texinfo](#), and [Sphinx](#).

Stable, transparent, and open software with few dependencies are easily reviewable. See section [3.1](#) for possible reviewing processes.

2.7.3 GROMACS

GROMACS [Abraham et al. 2015; Abraham et al. 2025] is a collection of programs for performing simulations of molecular systems, with a focus on biological macromolecules and on very large simulations requiring high-performance computing (HPC) resources. It is a typical example of domain-specific scientific software: well-known among practitioners of biomolecular simulations, but almost unknown outside of this community.

Scope: GROMACS is neither wide-spectrum nor situated, but somewhere in between.

Openness: GROMACS is Open Source software. It is a small ecosystem by itself since it consists of multiple programs designed to work together. This ecosystem is, however, not convivial. The source code is far too complex for the typical GROMACS user to understand, and some of the individual tools are opaque. Non-conviviality is confirmed by the observation that the developer community is a small subset of the user community.

Stability: GROMACS development follows a well-documented release cycle, with one major release per year and several bug-fix releases over a period of two years. This means that at any given time, two major releases are supported by bug fixes. Major releases mostly add functionality, but also

improve performance and add support for new computing systems, which is an important aspect for HPC where performance requires adapting to specific processors, memory layouts, and communication networks. The basic functionality is intended to be stable, but its implementation is not because it evolves due to ongoing optimization and porting efforts. Whether or not the basic functionality actually *is* stable is hard to decide because of opacity (see below). GROMACS is therefore evolving software.

Transparency: The GROMACS code is a highly optimized implementation of several computational models for molecular interactions (called *force fields*) and several simulation methods. These ingredients are inextricably interwoven in the code because of the primacy of performance. The implemented models and methods are described and discussed in the scientific literature, and also implemented in other software. However, no precise and complete description exists anywhere outside of software source code. There is no point in even asking if GROMACS implements any such model or method correctly, as there is no precise notion of correctness. For this reason, GROMACS is highly opaque.

Dependencies: The core dependencies are a C/C++ compiler, the [FFTW](#) library for computing Fourier transforms, an implementation of the [BLAS](#) and [LAPACK](#) specifications for linear algebra operations, plus the build system [cmake](#). All of these packages, as well as their respective dependencies, are well-known stable elements of computational infrastructure. However, GROMACS depends critically on properties of specific compilers on specific machines, due to its high-performance focus. It also has additional dependencies for specific types of computers, which include proprietary libraries for using nVIDIA GPUs. A quick glance at the [GROMACS installation instructions](#) gives an idea of the potential complexity of its dependency graph.

GROMACS is a good example for scientific software that is impossible to review in its entirety, due to the combination of being opaque, evolving, and having complex and platform-dependent dependencies. At best, a review could cover a specific version running on a specific platform used for a narrowly defined type of application. Opacity and instability also increase the risk of inappropriate use by scientists who have an insufficient or outdated mental model of the software’s operation. I have personally witnessed several cases in which inappropriate use of GROMACS led to obviously wrong results, which suggests that there are also cases that go unnoticed.

Given today’s development tools and best practices in software engineering, plus the focus on performance, it is impossible to make the software more reviewable. For potential measures to design more reviewable molecular simulation software, see sections [3.3](#), [3.4](#) and [3.5](#).

3 Improving the reviewability of automated reasoning systems

The analysis presented in the previous section can by itself improve the basis for trust in automated reasoning, by providing a vocabulary for discussing reviewability issues. Ensuring that both developers and users of scientific software are aware of where the software is located on the different scales I have described makes much of today’s tacit knowledge about scientific software explicit, avoiding misplaced expectations.

However, scientists can also work towards improving their computational practices in view of more reviewable and thus ultimately more reliable results. These improvements include both new reviewing processes, supported by institutions that remain to be created, and new software engineering practices that take into account the specific roles of software in science, which differ in some important respects from the needs of the software industry. The four measures I will explain in the following are summarized in Fig. 3.

Review the reviewable
Emphasize situated and convivial software
Make scientific software explainable
Use Digital Scientific Notations

Figure 3: Four measures that can be taken to make scientific software more trustworthy.

3.1 Review the reviewable

As my analysis has shown, some types of scientific software are reviewable, but not reviewed today. Several scientific journals encourage authors to submit code along with their articles, but only a small number of very specialized journals (e.g., [Computo](#), the [Journal of Digital History](#), [ReScience C](#)) actually review the submitted code, which tends to be highly situated. Other journals, first and foremost the [Journal of Open Source Software](#), review software according to generally applicable criteria of usability and software engineering practices, but do not expect reviewers to judge the correctness of the software nor the accuracy or completeness of its documentation. This would indeed be unrealistic in the standard journal reviewing process that asks a small number of individual researchers to evaluate, as volunteers and within short delays, submissions that are often only roughly in their field of expertise and represent the work of large teams over many years.

The first category of software that is reviewable but not yet reviewed is stable wide-spectrum software, such as the GNU Scientific Library (see section 2.7.2). Reviewing could take the form of regular audits, performed by experts working for an institution dedicated to this task. In view of the wide use of the software by non-experts in its domain, the audit should also inspect the software’s documentation, which needs to be up to date and explain the software’s functionality with all the detail that a user must understand. Specifications would be particularly valuable in this scenario, as the main interface between developers, users, and auditing experts. For opaque software, formal specifications could even be made a requirement, in the interest of an efficient audit. The main difficulty in achieving such audits is that none of today’s scientific institutions consider them part of their mission.

The second category of reviewable software is transparent software, for which reviewing consists of mostly the same steps as the established practices of verification and validation: examining the output for well-known inputs and/or comparing to the outputs of alternative implementations of the same tasks. The missing step is the publication of reports about verification and validation processes performed by recognized independent experts.

The third category of reviewable software contains situated software, which can and should be reviewed together with the other outputs of a research project. For small projects, in terms of the number of co-authors and the disciplinary spread, situated software could be reviewed as part of today’s peer review process, managed by scientific journals. The experience of pioneering journals in this activity could be the basis for elaborating more widely applied reviewing guidelines. For larger or multidisciplinary projects, the main issue is that today’s peer review process is not adequate at all, even in the (hypothetical) complete absence of software. Reviewing research performed by a multidisciplinary team requires another multidisciplinary team, rather than a few individuals reviewing independently. The integration of situated software into the process could provide the occasion for a more general revision of the peer review process.

3.2 Science vs. the software industry

In the first decades of computing technology, scientific computing was one of its main application domains, alongside elaborate bookkeeping tasks in commerce, finance, and government. Many computers (e.g. Digital Equipment Corporation’s PDP and VAX series), operating systems (e.g. International Business Machine’s VM/CMS), and compilers (e.g. Fortran) were designed for the needs of scientists. Today, scientists use mostly computers designed and produced for personal and business needs. Even supercomputers are constructed to a large degree from high-grade commodity components that

are also found in servers for cloud computing. Much infrastructure software, such as operating systems or compilers, are also commodity products developed primarily for other application domains.

From the perspective of development costs, this evolution makes economic sense. However, as with any shift towards fewer but more general products serving a wider client base, the needs of the larger client groups take priority over those of the smaller ones. Unfortunately for science, it is today a relative small application domain for software technology.

In terms of my analysis of reviewability in section 2, the software industry has a strong focus on proprietary wide-spectrum software, with a clear distinction between developers and users. Opacity for users is not seen as a problem, and may even be considered advantageous as a barrier to reverse-engineering of the software by competitors. Stability is an expensive characteristic that only relatively few customers (e.g. banks or some branches of industry) are willing to pay for. In contrast, novelty is an important selling argument in many profitable application domains, leading to attitudes such as “move fast and break things” (the long-time motto of Facebook founder Mark Zuckerberg), and thus favoring evolving software.

As a consequence of the enormous growth of non-scientific compared to scientific software, today’s dominant software development tools and software engineering practices largely ignore situated and convivial software, the impact of dependencies, and the scientific method’s requirement for transparency. However, it can be expected that the ongoing establishment of Research Software Engineers as a specialization at the interface between scientific research and software engineering will lead to development practices that are better aligned with the specific needs of science. It is such practices that I will propose in the following sections.

3.3 Emphasize situated and convivial software

As mentioned in section 2.1, many important scientific software packages are domain-specific tools and libraries, which have neither the large user base of wide-spectrum software that justifies external audits, nor the narrow focus of situated software that allows for a low-effort one-time review by domain experts. Developing suitable intermediate processes and institutions for reviewing such software is perhaps possible, but I consider it scientifically more appropriate to restructure such software into a convivial collection of more situated modules, supported by a shared layer that is reviewable due to being wide-spectrum or transparent. However, this implies assigning a lower priority to reusability, in conflict with today’s best practices in software engineering, and with recent initiatives to apply the FAIR principles to software [Barker et al. 2022].

In such a scenario, a domain library becomes a collection of source code files that implement core models and methods, plus ample documentation of both the methods and implementation techniques. The well-known book “Numerical Recipes” [Press et al. 2007] is a good example for this approach. Code, documentation, test, and examples need to be written together, with the explicit goal of supporting users in understanding and modifying the code, in the spirit of literate programming [Knuth 1984].

Users whose needs are not fully met by such a library would make a copy of the source code files relevant for their work, adapt them to the particularities of their applications, and make them an integral part of their own project. In terms of FLOSS jargon, users would make a partial fork of the project. Version control systems of a kind that doesn’t exist yet would ensure provenance tracking and support the discovery of other forks. Keeping up to date with relevant forks of one’s software would become part of everyday research work, much like keeping up to date with publications in one’s wider research community. In fact, another way to describe this approach is full integration of scientific software development into established research practices, rather than keeping it a distinct activity governed by different rules. Yet another perspective is giving priority to the software’s role as a representation of scientific knowledge over its role as a tool [Hinsen 2014].

Since this approach differs radically from anything that has been tried in practice so far, it is premature to discuss its advantages and downsides. Only practical experience can show to what extent pre-digital and pre-industrial forms of collaborative knowledge work can be adapted to automated reasoning. Nevertheless, I will indulge in some speculation on this topic, to give an idea of what we can fear or hope for.

On the benefit side, the code supporting a specific research project becomes much smaller and more understandable, mitigating opacity. Its computational environment is smaller as well, and entirely composed of stable wide-spectrum or transparent software. Reviewability is therefore much improved. Moreover, users are encouraged to engage more intensely with the software, reducing the risk of inappropriate use. The lower entry barrier to appropriating the code makes it accessible to a wider range of researchers, increasing inclusiveness and epistemic diversity.

The main loss I expect is in the efficiency of implementing and deploying new ideas. A strongly coordinated development team whose members specialize on specific tasks is likely to advance more quickly in a well-defined direction. This can be a disadvantage in particular for software whose complexity is dominated by technical rather than scientific aspects, e.g. in high-performance computing or large-scale machine learning applications.

3.4 Make scientific software explainable

Opacity is a major obstacle to the reviewability of software and results obtained with the help of software, as I have explained in section 2.3. Depending on one’s precise definition of opacity, it may be impossible to reduce it. Pragmatically, however, opacity can be mitigated by *explaining* what the software does, and providing tools that allow a scientist to *inspect* intermediate or final results of a computation.

The popularity of computational notebooks, which can be seen as scripts with attached explanations and results, shows that scientists are indeed keen on making their work less opaque. But notebooks are limited to the most situated top layer of a scientific software stack. Code cells in notebooks refer to library code that can be arbitrarily opaque, difficult to access, and to which no explanations can be attached.

An interesting line of research in software engineering is exploring possibilities to make complete software systems explainable [Nierstrasz and Girba 2022; Nierstrasz and Gîrba 2024]. The motivation is similar to explainable AI [Gunning et al. 2019], but the methods are very different, considering that the target of the explanation is human-authored source code. A new development environment, [Glamorous Toolkit](#) [feenk.com 2023], has been designed explicitly to explore these ideas. Although this work is motivated by situated business applications, the basic ideas should be transferable to scientific computing. The screenshots in [Nierstrasz and Gîrba 2024] convey a first impression of how these techniques work in practice. However, I recommend readers to actually download Glamorous Toolkit and explore it interactively.

The approach is based on three principles. The first one is the same as for computational notebooks: the integration of code with explanatory narratives that also contain example code and computed results. Unlike traditional notebooks, Glamorous Toolkit allows multiple narratives to reference a shared codebase of arbitrary structure and complexity. The second principle is the generous use of examples, which serve both as an illustration for the correct use of the code and as test cases. In Glamorous Toolkit, whenever you look at some code, you can access corresponding examples (and also other references to the code) with a few mouse clicks. The third principle is what the authors call *moldable inspectors*: situated views on data that present the data from a domain perspective rather than in terms of its implementation. These three techniques can be used by software developers to facilitate the exploration of their systems by others, but they also support the development process itself by creating new feedback loops.

3.5 Use Digital Scientific Notations

As I have briefly mentioned in the introduction, specifications are contracts between software developers and software users that describe the expected behaviour of the software. Formal specifications are specifications written in a formal language, i.e. a language amenable to automated processing. There are various techniques for ensuring or verifying that a piece of software conforms to a formal specification [Wikipedia 2025a]. The use of these tools is, for now, reserved to software that is critical for safety or security, because of the high cost of developing specifications and using them to verify implementations, although the idea of *lightweight* formal methods, with the potential of much wider applicability, has been circulating for almost thirty years [Jackson and Wing 1996; Zamansky et al. 2018].

Technically, formal specifications are *constraints* on algorithms and programs, in much the same way as mathematical equations are constraints on mathematical functions [Hinsen 2023]. Such constraints are often much simpler than the algorithms they define. As an example, consider the task of sorting a list. The (informal) specification of this task is: produce a new list whose elements are (1) the same as those of the input list and (2) sorted. A formal version requires some additional details, in particular a definition of what it means for two lists to have “the same” elements, given that elements can appear more than once in a list. There are many possible algorithms conforming to this specification, including well-known sorting algorithms such as quicksort or bubble sort [Wikipedia 2025b]. All of them are much more elaborate than the specification of the result they produce. As a consequence, they are much more difficult to understand. Consider an implementation of, say, Quicksort. Many such implementations are available on [Rosetta Code](#). It is not obvious from reading such code that it sorts a list. It is even less obvious that it does so correctly, as there are many details where a small change would make the results incorrect without this being obvious. The specification, on the other hand, is immediately understandable. Moreover, specifications are usually more modular than algorithms, which also helps human readers to better understand what the software does [Hinsen 2023].

The software engineering contexts in which formal specifications are used today are very different from the potential applications in scientific computing that I outline here. In software engineering, specifications are written to formalize the expected behavior of the software *before* it is written. The software is considered correct if it conforms to the specification. In scientific research, software evolves in parallel with the scientific knowledge that it encodes or helps to produce. A formal specification has to evolve in the same way, and is best seen as the formalization of the scientific knowledge. Change can flow from specification to software, but also in the opposite direction. Moreover, most specifications are likely to be incomplete, leaving

out aspects of software behavior that are irrelevant from the point of view of science (e.g. resource management or technical interfaces such as Web APIs), but also aspects that are still under exploration and thus not yet formalized. For these reasons, I prefer the term *Digital Scientific Notation* [Hinsen 2018b], which better expresses the role of formal specifications in this context.

Digital Scientific Notations can take many forms. They do not have to resemble programming languages, nor the specification languages used in software engineering. My own experimental Digital Scientific Notation, *Leibniz* [Hinsen 2024a], is intended to resemble traditional mathematical notation as used e.g. in physics. Its statements are embeddable into a narrative, such as a journal article, and it intentionally lacks typical programming language features such as scopes that do not exist in natural language, nor in mathematical notation. For a simple example, see [the Lotka-Volterra equations in Leibniz](#). For other domains, graphical notations may be more appropriate. These notations, the tooling that integrates them with software, and the scientific practices for working with them, all remain to be developed. The main expected benefits are conviviality of the specifications and transparency of the tools that process them.

4 Conclusion

My principal goal with this work is to encourage scientists and research software engineers to reflect about their computational practices. Why, and to what degree, do you trust your own computations? What are your reliability indicators? Do you expose them in your publications? How reliable does your software have to be to support the conclusions you draw from their results? Why, and to what degree, do you trust the computations in the papers you read and cite? Do you consider their reliability sufficient to support the conclusions made?

These questions are abstract. Answering them requires considering the concrete level of the specific software used in a computation. The five categories I have discussed in section 2 should help with this step, even though it may be difficult at first to evaluate the software you use on some of the scales. The case studies of section 2.7 should help. Situated software is easy to recognize by its narrow application domain. The size of a software environment is not difficult to measure, but it requires appropriate tools and training in their use. The evaluation of stability is often not difficult, but requires a significant effort, in particular an examination of a software project's history. Conviviality is hard to diagnose, but rare anyway. This reduces the examination to Open Source vs. proprietary, which is straightforward.

This leaves the transparency vs. opacity scale, which deserves a more detailed discussion. Most experienced computational scientists make sure to examine both intermediate and final results for plausibility, making use of known properties such as positivity or order of magnitude. But plausibility is a fuzzy concept. Software is transparent only if users can check results for *correctness*, not mere plausibility. The strategies I proposed (sections 3.3, 3.4 and 3.5) have the goal of making such correctness checks easier. If plausibility is all we can check for, then the software is opaque, and its users are faced with a dilemma when their results are neither obviously correct nor obviously wrong: are they entitled to consider them good enough? In practice they do, because the only realistic alternative would be to stop using computers. Soergel [2015] and Thimbleby [2023] consider this “trust by default” misplaced, given what software engineering research tells us about the frequency of mistakes, and I agree. The examples from the reproducibility crisis (see the Introduction) support this view that scientists tend to overestimate the reliability of their work in the absence of clear signs of problems.

From the perspective of computational reliabilism, the two most critical characteristics are stability and transparency. If either one is lacking, it is difficult to consider software to be reliable. Evolving software changes too frequently to permit the accumulation of evidence for reliability. Opaque software does not have a clearly defined behavior. If you don’t know what a piece of software is supposed to do exactly, it is hard to affirm that it does so reliably.

The ideal structure for a reviewable scientific software stack (see Fig. 2) whose reliability can actually be argued for would consist of a foundation of stable infrastructure and a top layer of situated software (a script, a notebook, or a workflow) that orchestrates the computations answering a specific scientific question. The intermediate domain-specific layer would have to be redesigned such that each of its components is either transparent or small. Digital scientific notations can help to achieve transparency. Situatedness is a good strategy for keeping code small.

The main issues we face today are evolving and opaque software, both of which are common in science. Reaching stability requires time, a large enough user base, and high software engineering standards. Mitigating opacity, e.g. by adopting the strategies I have proposed, requires a significant effort. Reliability comes at a cost. Making good choices requires a cost-benefit analysis in the context of each specific research project. The arguments for the choice should be mentioned in every research report, to permit readers an assessment of the reliability of the reported findings.

The difficulty of reviewing scientific software also illustrates the deficiencies

of the current digital infrastructure for science.⁷ The design, implementation, and maintenance of such an infrastructure, encompassing hardware, software, and best practices, has been neglected by research institutions, in spite of an overt enthusiasm about the scientific progress made possible by digital technology. The situation is improving for research data, for which appropriate repositories and archives are becoming available. For software, the task is more complex, and hindered by the contagious neophilia (“tech churn”) of the software industry. Scientists, research software engineers, research institutions, and funding agencies must recognize the importance of stable and reliable infrastructure software, which requires long-term funding and inclusive governance.

Beyond infrastructural issues, reviewing code is considered impractical by many researchers because of the enormous effort it seems to require. However, this argument ignores the fact that we have not yet made a significant effort to create appropriate reviewing processes. It also ignores the new possibilities opened up by large language models for assisting with reviewing code, and in particular reviewing code and its documentation together as a single entity. But perhaps most importantly, the argument ignores second-order effects of reviewing. In the long run, the practice of reviewing scientific code will create an incentive for scientists and engineers to make their code more reviewable. And the best time to start the development of reviewing practices is right now. The use of large language models in software development is growing. It can help to make software more reviewable, but also less reviewable. The best way to push the lever to “more reviewable” is starting to review now.

Finally, I need to address a frequent objection to any criticism of the current use of automated reasoning in science. The objection is that there is no problem, that science works pretty well overall, because potential mistakes in an individual scientific study tend to get detected in a later stage of the eternal knowledge refinement process of research. While this is true in principle, it is important to understand the cost of postponing the detection of mistakes and other issues (biases, dubious approximations, etc.) to later stages. Consider the case of the five retracted protein structure papers [Miller 2006] that I mentioned in the Introduction. They represent a significant amount of work, much of which was wasted. In between their publication and their retraction, other researchers working on similar proteins had their work rejected because it was in contradiction with these high-profile publications, meaning more wasted effort. The undetected software bug caused a lot of frustration, probably ruined the careers of a few young researchers, and wasted scarce research money. On the other hand, independent critical inspection of the software would probably have prevented all of that.⁸

⁷For more examples, see [Saunders 2022].

⁸Assuming that the explanation given by the researchers for their mistake is correct.

The question is thus not if we *must* improve computational practices, but if we *should* do so, in order to reduce frustration and hardship to researchers, monetary loss to science funders, and lost opportunities to society at large.

5 Acknowledgments

I would like to thank Nico Formanek [Formanek 2025] and an anonymous reviewer [Anonymous Reviewer 2025] for helpful reviews of an earlier version of this article, and Juan Durán for an insightful discussion about computational reliabilism.

References

- ABRAHAM, M., ALEKSEENKO, A., ANDREWS, B., ET AL. 2025. [GROMACS 2025.3 Manual](#).
- ABRAHAM, M.J., MURTOLA, T., SCHULZ, R., ET AL. 2015. [GROMACS: High performance molecular simulations through multi-level parallelism from laptops to supercomputers](#). *SoftwareX* 1–2, 19–25.
- ANONYMOUS REVIEWER. 2025. [Review: Establishing trust in automated reasoning \(Round 1 - Review 1\)](#).
- BAKER, M. 2016. [1,500 scientists lift the lid on reproducibility](#). *Nature* 533, 7604, 452–454.
- BARKER, M., CHUE HONG, N.P., KATZ, D.S., ET AL. 2022. [Introducing the FAIR Principles for research software](#). *Scientific Data* 9, 1, 622.
- BILKE, L., FISCHER, T., NAUMOV, D., AND MEISEL, T. 2025. [Reproducible HPC software deployments, simulations, and workflows – a case study for far-field deep geological repository assessment](#). *Environmental Earth Sciences* 84, 17, 502.
- CHAPTER 1: FOR WANT OF A PRINTER. 2002. In: *Free as in freedom: Richard Stallman’s crusade for free software*. O’Reilly, Sebastopol, Calif. : Farnham.
- COLLIARD, J.-E., HURLIN, C., AND PÉRIGNON, C. 2022. [The Economics of Computational Reproducibility](#).
- COURTÈS, L. 2021. What’s in a package. *GuixHPC blog*.
- DUBOIS, P.F., HINSEN, K., AND HUGUNIN, J. 1996. [Numerical Python](#). *Computers in Physics* 10, 3, 262.
- DURÁN, J.M. 2025. Beyond transparency: Computational reliabilism as an externalist epistemology of algorithms. <https://arxiv.org/abs/2502.20402>.
- DURÁN, J.M. AND FORMANEK, N. 2018. [Grounds for Trust: Essential Epistemic Opacity and Computational Reliabilism](#). *Minds and Machines* 28, 4, 645–666.
- FEENK.COM. 2023. Glamorous Toolkit.

The supposedly faulty code has never been made public.

- FORMANEK, N. 2025. [Review: Establishing trust in automated reasoning \(Round 1 - Review 2\)](#).
- FRENCH MINISTRY FOR HIGHER EDUCATION, RESEARCH, AND INNOVATION. 2021. [Second French Plan-for-Open-Science web](#).
- GALASSI, M., ED. 2009. *GNU scientific library reference manual: For GSL version 1.12*. Network Theory, Bristol.
- GEORGIU, P., ZANOS, P., MOU, T.-C.M., ET AL. 2022. [Experimenters' sex modulates mouse behaviors and neural responses to ketamine via corticotropin releasing factor](#). *Nature Neuroscience* 25, 9, 1191–1200.
- GONZALEZ-BARAHONA, J.M. 2021. [A Brief History of Free, Open Source Software and Its Communities](#). *Computer* 54, 2, 75–79.
- GSL - GNU SCIENTIFIC LIBRARY - GNU PROJECT - FREE SOFTWARE FOUNDATION. 2025.
- GUNNING, D., STEFIK, M., CHOI, J., MILLER, T., STUMPF, S., AND YANG, G.-Z. 2019. [XAI—Explainable artificial intelligence](#). *Science Robotics* 4, 37, eaay7120.
- HARRIS, C.R., MILLMAN, K.J., VAN DER WALT, S.J., ET AL. 2020. [Array programming with NumPy](#). *Nature* 585, 7825, 357–362.
- HASSELBRING, W., DRUSKAT, S., BERNOTH, J., ET AL. 2025. [Multidimensional Research Software Categorization](#). *Computing in Science & Engineering* 27, 2, 59–68.
- HERNDON, T., ASH, M., AND POLLIN, R. 2014. [Does high public debt consistently stifle economic growth? A critique of Reinhart and Rogoff](#). *Cambridge Journal of Economics* 38, 2, 257–279.
- HETTRICK, S. 2016. *Research Software Sustainability*. Knowledge Exchange.
- HETTRICK, S., ANTONIOLETTI, M., CARR, L., ET AL. 2014. [UK Research Software Survey 2014](#).
- HINSEN, K. 2014. [Computational science: Shifting the focus from tools to models](#). *F1000Research* 3, 101.
- HINSEN, K. 2016. [The Power to Create Chaos](#). *Computing in Science & Engineering* 18, 4, 75–79.
- HINSEN, K. 2018a. [Reusable Versus Re-editable Code](#). *Computing in Science & Engineering* 20, 3, 78–83.
- HINSEN, K. 2018b. [Verifiability in computer-aided research: The role of digital scientific notations at the human-computer interface](#). *PeerJ Computer Science* 4, e158.
- HINSEN, K. 2019. [Dealing With Software Collapse](#). *Computing in Science & Engineering* 21, 3, 104–108.
- HINSEN, K. 2020. Computational Reproducibility. In: *Computation in Science: From concepts to practice*. IOP Publishing.
- HINSEN, K. 2023. [The nature of computational models](#). *Computing In Science & Engineering* 25, 1, 61–66.
- HINSEN, K. 2024a. [Leibniz - a Digital Scientific Notation](#).
- HINSEN, K. 2024b. [Redressing the Balance: A Yin-Yang Perspective on](#)

- Information Technology. *Proceedings of the 2024 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, Association for Computing Machinery, 194–204.
- HINSEN, K. 2025. [Reviewing Research Software](#). *Computing in Science & Engineering* 27, 3, 64–66.
- HOCQUET, A. AND WIEBER, F. 2017. [“Only the Initiates Will Have the Secrets Revealed”: Computational Chemists and the Openness of Scientific Software](#). *IEEE Annals of the History of Computing* 39, 4, 40–58.
- HUMPHREYS, P. 2009. [The Philosophical Novelty of Computer Simulation Methods](#). *Synthese* 169, 3, 615–626.
- ILLICH, I. 1973. *Tools for conviviality*. Calders and Boyars, London.
- JACKSON, D. AND WING, J.M. 1996. Lightweight Formal Methods. *IEEE Computer April 1996*, 21–22.
- JUMPER, J., EVANS, R., PRITZEL, A., ET AL. 2021. [Highly accurate protein structure prediction with AlphaFold](#). *Nature* 596, 7873, 583–589.
- KAY, A.C. 1972. [A Personal Computer for Children of All Ages](#). *Proceedings of the ACM annual conference - Volume 1*, Association for Computing Machinery.
- KELL, S. 2020. [Convivial design heuristics for software systems](#). *Conference Companion of the 4th International Conference on Art, Science, and Engineering of Programming*, ACM, 144–148.
- KNUTH, D.E. 1984. Literate programming. *The Computer Journal* 27, 2, 97–111.
- KO, A.J., ABRAHAM, R., BECKWITH, L., ET AL. 2011. [The state of the art in end-user software engineering](#). *ACM Computing Surveys* 43, 3, 21:1–21:44.
- KORTZFLEISCH, V.T. VON, AMBRÉE, O., KARP, N.A., ET AL. 2022. [Do multiple experimenters improve the reproducibility of animal studies?](#) *PLOS Biology* 20, 5, e3001564.
- LEONELLI, S. 2022. [Open Science and Epistemic Diversity: Friends or Foes?](#) *Philosophy of Science* 89, 5, 991–1001.
- LONGO, L., BRCIC, M., CABITZA, F., ET AL. 2024. [Explainable Artificial Intelligence \(XAI\) 2.0: A manifesto of open challenges and interdisciplinary research directions](#). *Information Fusion* 106, 102301.
- MALKA, J., ZACCHIROLI, S., AND ZIMMERMANN, T. 2024. Reproducibility of Build Environments through Space and Time. *46th International Conference on Software Engineering (ICSE 2024) - New Ideas and Emerging Results (NIER) Track*.
- MERALI, Z. 2010. [Computational science: ...Error](#). *Nature* 467, 7317, 775–777.
- MILLER, G. 2006. [A Scientist’s Nightmare: Software Problem Leads to Five Retractions](#). *Science* 314, 5807, 1856–1857.

- NANGIA, U. AND KATZ, D.S. 2017. [Understanding Software in Research: Initial Results from Examining Nature and a Call for Collaboration](#). *2017 IEEE 13th International Conference on e-Science (e-Science)*, 486–487.
- NARDI, B.A. 1993. *A small matter of programming: Perspectives on end user computing*. MIT Press, Cambridge, MA.
- NIELSEN, M. 2023. How is AI impacting science?
- NIERSTRASZ, O. AND GIRBA, T. 2022. [Making Systems Explainable](#). *2022 Working Conference on Software Visualization (VISSOFT)*, IEEE, 1–4.
- NIERSTRASZ, O.M. AND GÎRBA, T. 2024. [Moldable Development Patterns](#). *Proceedings of the 29th European Conference on Pattern Languages of Programs, People, and Practices*, Association for Computing Machinery, 1–14.
- PARASURAMAN, R. AND RILEY, V. 1997. [Humans and Automation: Use, Misuse, Disuse, Abuse](#). *Human Factors* 39, 2, 230–253.
- POTE, T. 2025. Stability by Design.
- PRESS, W.H., TEUKOLSKY, S.A., VETTERLING, W.T., AND FLANNERY, B.P. 2007. *Numerical recipes: The art of scientific computing*. Cambridge University Press, Cambridge, UK ; New York.
- SAMUEL, S. AND MIETCHEN, D. 2024. [Computational reproducibility of Jupyter notebooks from biomedical publications](#). *GigaScience* 13, giad113.
- SAUNDERS, J.L. 2022. Decentralized Infrastructure for (Neuro)science. <https://arxiv.org/abs/2209.07493>.
- SHAW, M. 2022. [Myths and mythconceptions: What does it mean to be a programming language, anyhow?](#) *Proceedings of the ACM on Programming Languages* 4, HOPL, 1–44.
- SHIRKY, C. 2004. Situated Software.
- SOERGEL, D.A.W. 2015. [Rampant software errors may undermine scientific results](#).
- SPIER, R. 2002. [The history of the peer-review process](#). *Trends in Biotechnology* 20, 8, 357–358.
- STODDEN, V. AND MIGUEZ, S. 2013. [Best Practices for Computational Science: Software Infrastructure and Environments for Reproducible and Extensible Research](#). *Social Science Research Network* 2322276.
- STODDEN, V., SEILER, J., AND MA, Z. 2018. [An empirical analysis of journal policy effectiveness for computational reproducibility](#). *Proceedings of the National Academy of Sciences* 115, 11, 2584–2589.
- THIMBLEBY, H. 2023. [Improving Science That Uses Code](#). *The Computer Journal*, bxad067.
- THOMPSON, K. 1984. [Reflections on trusting trust](#). *Communications of the ACM* 27, 8, 761–763.
- TURKLE, S. 2009. *Simulation and its discontents*. The MIT Press, Cambridge, Massachusetts.
- WANG, J., LI, L., AND ZELLER, A. 2021. [Restoring Execution Environments](#)

- of Jupyter Notebooks. *2021 IEEE/ACM 43rd International Conference on Software Engineering (ICSE)*, 1622–1633.
- WIKIPEDIA. 2025a. Formal methods. *Wikipedia*.
- WIKIPEDIA. 2025b. Sorting algorithm. *Wikipedia*.
- WILKINSON, M.D., DUMONTIER, M., AALBERSBERG, IJ.J., ET AL. 2016. [The FAIR Guiding Principles for scientific data management and stewardship](#). *Scientific Data* 3, 160018.
- ZAMANSKY, A., SPICHKOVA, M., RODRIGUEZ-NAVAS, G., HERRMANN, P., AND BLECH, J.O. 2018. [Towards Classification of Lightweight Formal Methods](#). *Proceedings of the 13th International Conference on Evaluation of Novel Approaches to Software Engineering*, SCITEPRESS - Science and Technology Publications, Lda, 305–313.