OCD As Aberrant Integration of Uncertainty from Local Goals to Global Goals

Andra Geana [1], Christina L. Boisseau [2], Steven Rasmussen [3,4], Michael J. Frank [5]

1. Providence College, Providence, RI; 2. Northwestern Feinberg School of Medicine, Northwestern, Chicago, IL 3. Brown University Alpert Medical School, Providence, RI. 4. Butler Hospital, Providence RI. 5. Brown University Carney Institute, Providence, RI

This paper is a preprint and has not yet been peer-reviewed or published

Abstract

Compulsivity is a key feature of obsessive-compulsive disorder (OCD), with compulsions—repetitive, stereotyped behaviors and mental acts meant to relieve distress—often posing significant obstacles to daily function. The last few years have seen increases in compulsive behaviors; yet, the specific mechanisms of compulsions in OCD remain unknown. Previous research holds mixed findings on compulsivity-linked cognitive deficits, and, while many accounts assume compulsions aim to reduce anxiety, most individuals with anxiety do not develop compulsions—suggesting that the anxiety-reduction account is incomplete. A deeper insight into the neural and cognitive mechanisms underpinning compulsions holds immediate potential benefits for improving treatment options for OCD patients. We propose here that one factor in the formation of compulsions is the inability to integrate uncertainty across hierarchical goal levels. Using a predictive inference task with a hierarchical structure relying on local uncertainty reduction (short-term goals) in the service of global uncertainty reduction (long-term goal), we tested local and global learning in OCD patients (n = 20) and healthy age-matched controls (n = 20). Both groups showed the ability to reduce uncertainty locally, learning the statistical structure of the local environment based on observed data; however, only the healthy volunteers showed evidence of integrated learned knowledge at the local level into reducing uncertainty about the higher level of the task hierarchy. We suggest that this indicates a potential mechanism for compulsions that relies on the inability to adaptively prioritize and "toggle" among different local goals in the service of global goals.

Keywords: obsessive-compulsive disorder; compulsivity; computational modelling; OCD; learning; reinforcement learning

Introduction

The prevalence of obsessive-compulsive disorder (OCD) has increased worldwide in the last few years since the COVID-19 pandemic[1], with contamination-related compulsions found at unprecedented levels even in previously undiagnosed populations[2]. Such compulsions, frequently manifesting as 'checking' or other elaborate routines to resolve uncertainty about the state of the world, force an individual to spend disproportionate time and effort engaged in a repetitive cycle, unable to break out and focus on other things. For example, while a healthy person might choose to wash their hands after returning home from a ride on public transit, then (once their hands are clean) move on with usual daily life, a person with OCD might fear contamination even after handwashing. This fear might lead them to wash their hands again after touching their clothes or touching a doorknob; then, perhaps, they might choose to clean the doorknob, or mop the floor to ensure they did not track in contaminants – and then wash their hands again and again after each step.

Yet, despite rising prevalence and severe impact on daily life—OCD ranks among the leading causes of disability in the United States, and in up to 30% of cases, treatment (including combinations of medication and therapy) fails to significantly relieve compulsions [3, 4]—a mechanistic understanding of the root of compulsions is largely lacking. This is due in part to the variability in symptoms and high comorbidity with other disorders, and to the fact that OCD is unlikely to result from alteration of any singular molecular mechanism. For instance, some accounts frame compulsions from a learning-failure perspective, assuming deficits in the ability to correctly represent action-outcome links [5, 33]. Others suggest a lack of goal-directed control, proposing that OCD impairs the ability to act appropriately in pursuit of goals and reverting to habits [5, 6]. Empirical data shows mixed evidence, with OCD individuals able to learn comparably to healthy volunteers on a variety of tasks, and goal-directedness failures inconsistent across tasks and domains [7-10]. Crucially, while anxiety is considered a vital factor in OCD, it is worth nothing that most people suffering from anxiety or intrusive anxious thoughts do not develop compulsions [11]; thus, it is possible that mechanisms independent of anxiety contribute to the formation of compulsions, and elucidating these separate mechanisms is critical to our understanding of the disorder.

Compulsivity as a Deficit in Uncertainty-Related Processes While classical neurocognitive assessments of OCD often involve diagnostic interviews, questionnaires (including clinical scales) and self-reported measures [12], observation from clinical settings suggests that many compulsive behaviors are aimed at reducing uncertainty, with the inability to do so ultimately perpetuating the irrational thoughts and behaviors [13]. Uncertainty plays a role in most learning and decision scenarios, from ordering dinner to choosing a job, and evidence across species shows that the ability to represent and reduce uncertainty is essential for optimal decision-making [14-15, 34-35]. Failures to regulate uncertainty have been linked to maladaptive learning in a variety of neuropsychiatric disorders [16-20]. Individuals with OCD in particular exhibit a profound intolerance of uncertainty , and perhaps as a result, oversample in uncertain environments [21-23]. However, uncertainty-driven computations remain relatively under-examined in the context of compulsivity, and it remains unclear whether patients simply show a heightened aversion to uncertainty *per se*. For example, patients might assign a more negative value to the prospect of the undesired outcome for a given level of uncertainty. Or, they might exhibit overall greater uncertainty due to disruptions in the computations needed to reduce it. Here, we consider the latter possibility.

It is possible that one source of mixed results in lab studies on learning deficits in OCD is the commonly-used task structure that assesses uncertainty about a single decision (e.g., learning the underlying reward structures of two independent options). In contrast, most real-world learning scenarios involve a hierarchy of goals, with “local” subgoals (e.g., “will washing my hands make them cleaner?”) in the service of broader, “global” goals (“how can I stay healthy?”). Under such framework, OCD patients may be able to resolve uncertainty about the outcomes of a local decision, but then fail to update the uncertainty about progress toward the global goal, due to disruptions in the mechanism responsible for integrating across timescales and hierarchical levels of abstraction. In line with this hypothesis, recent work [22, 24-25] shows impaired transfer of information (in the form of explicit and implicit task-relevant memories) across repeated decisions. However, to our knowledge, the integration of information in the service of short-term and long-term goals has never been tested in OCD; a mismatch in the integration of these parallel learning processes into global goals can produce compulsive behaviors.

Using a predictive inference task with a hierarchical learning structure, we tested this hypothesis in a sample of individuals with OCD and age-matched controls. While OCD patients showed intact local learning of stochastic outcomes, they exhibited impaired ability to leverage this learning to update their beliefs about global structure. We propose that a mismatch in the integration of these local learning processes into global goals can produce compulsive behaviors, while preserving the ability to locally reduce uncertainty.

Methods

Participants

Participants were recruited through the OCD Research Program at Butler Hospital. Out of the 463 participants that were phone screened for initial eligibility, 55 (8.4%) met criteria for in-person evaluation for final determination of eligibility. To be included in the OCD group participants were required to be adults age 18-55, meet DSM-5 criteria for OCD, and identify OCD as their primary psychiatric concern. Exclusion criteria for the OCD group included past month substance use disorder or report of lifetime bipolar or psychotic disorders. Healthy controls were free of current psychiatric disorders.

Of the 55 individuals that past the initial phone screen, 4 were ruled out of the OCD group because of subclinical symptoms, 3 were ruled out of the healthy control group due to the presence of psychopathology or unreliable report and an additional 5 did not show up to the in-person assessment. Thus, the final sample contained 20 participants with OCD and 20 healthy controls. The groups were matched on age and general intellectual ability as measured by the National Adult Reading Test-Revised (NART-R; Blair & Spreen,1989). All participants gave informed, written consent approved by the Institutional Review Board. They were compensated $80 dollars for approximately four hours of their time.

Measures

Psychiatric diagnoses were made by a trained postdoctoral fellow based on DSM-5 criteria using the Structured Clinical Interview for DSM-5 (SCID-5; First et al., 2015). Severity of obsessive and compulsive symptoms was assessed using the rater-administered Yale-Brown Obsessive-Compulsive Scale (YBOCS) and Symptom Checklist (Goodman et al., 1989a, b). Depression and anxiety symptoms were measured using the self-report Beck Depression Inventory (Beck et al., 1996) and Beck Anxiety Inventory (Beck et al., 1998). Finally, intolerance of uncertainty was measured using the self-report Intolerance of Uncertainty Scale (IUS, Buhr & Dugas, 2002).

The Archer Task

We developed a new predictive inference task, building upon the basic structure of previous tasks designed to challenge dynamic learning in uncertain (stochastic and volatile) environments [26; 27]. We modified this task (Fig 1) to include a hierarchical structure, with local outcomes generated through an overarching global structure. Players had to aim arrows to shoot enemies, the locations of which are hidden, but can be inferred from sequential observations. The task was framed as a game in which players must defeat an "evil overlord" by destroying his spaceship (located at an unknown position on the screen). The spaceship sends “cannons” to attack the archer, with each cannon launching up to thirty “minions” (*local samples*), sequentially. To introduce uncertainty, each minion’s position is drawn from a Gaussian distribution, whose *local mean* is centered on the position of the cannon. Each cannon’s position is, in turn, launched by the spaceship with stochastic variations around the *global mean*. This hierarchical structure ensured that observing each minion’s position can reduce uncertainty simultaneously on two levels. At the *local* level, the position of an individual minion provides a sample of evidence about the position of the current cannon, allowing the player to adjust their subsequent aim. With more samples they should have a more certain estimate about the cannon position (the local mean) and thus progressively adjust less with each stochastic outcome. At the *global* level, the player can use their updated estimate of the current cannon’s position to infer spaceship position.

The players earned a small reward (1 point) for accurately aiming arrows to hit each minion’s position, and a large reward at the end of the game (60 points) for accurately firing an arrow at the spaceship (they were instructed about this payoff structure during pre-task instructions as well as during the training games). At the end of the task, earned points were converted to a “bonus” reward of up to $5.

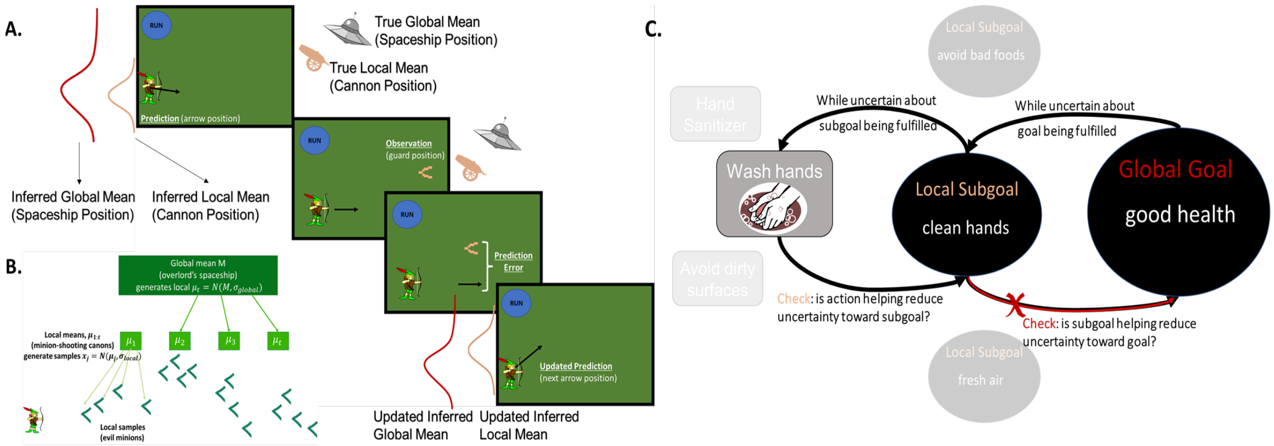

Figure 1: The Archer Task. ***A.*** *Example trial, starting with* ***prediction*** *(player's aim with the arrow), followed by the* ***observation*** *of a data point (minion) from the local distribution (canon), and the* ***updated prediction*** *(adjusting aim for next trial). The hierarchical structure of the task, with the two distributions governing minions and canon positions are show on the side of the screen. Panel* ***B*** *shows a graphical representation of the hierarchical Gaussian distribution, with the global distribution centered on a global mean (M) governing the distribution of local means, which in turn determine the position of each minion the player observers.* ***C.*** *Conceptual representation of how uncertainty can be used in the service of integration of local into global goals, and how the inability to integrate (bottom right-side arrow, red), even accompanied by intact ability to reduce uncertainty (left-side circuit) can lead to becoming "stuck" in local repetitive-action cycles.*

Trading off local and global reward A vital aspect to the task was that players had a limited number of 210 arrows. They were made aware of this limit at the beginning of the task, and during training; however, they did not see the remaining number arrows once training was completed and the task began. When all but one arrows were used up, the player would immediately encounter the spaceship and have one chance to fire at it. At that point, a more accurate estimate of the spaceship's location would translate into a better chance of earning the significantly higher global reward. On the other hand, each local canon could fire up to 30 minions. By sampling a series of minions, the player can be relatively confident about the canon's location and then accumulate local rewards (r = 1 point each) relatively easily (i.e., by aiming at the mean of their previous observations). To reflect this trade-off, on each trial, players were given the chance to decide whether to "stay" and face more minions from that cannon (thus improving their local estimate) or "run" and see a new cannon (i.e., sample more from the global structure and improve their global estimate).

In sum, to maximize total reward for the task, players had to learn both each local mean (i.e. learn where each cannon is, to accurately predict where the minions will come from) and the global mean (i.e. spaceship location) that generated all these local means. The speed of their learning on both these processes changed the optimal point for continuing or quitting each local game (see modelling below for formalizing this tradeoff).

All participants played the same number of trials (210), but due to the player-made decision to stay or leave, the number of games (waves of local minions) ranged from a minimum of 7 to a maximum of 15, with an average of just under 10 games.

Computational Modelling

For a deeper quantitative and mechanistic understanding, we take a computational psychiatry approach [28], mixing theory-driven and data-driven methods to predict behavior based on individual inferences about the structure and likelihood of different outcomes in the world. Assuming a hierarchy of long-term (global) and short-term (local) learning goals, our computational model of the archer task allows us to test to what degree participants performed parallel updating in the estimation of local and global uncertainty, and how they integrated the two decision problems to solve the task.

Uncertainty-Updating Model

On each trial (*t*) of each game (*g*), the model uses observed minion positions (*x*) to estimate the position of that game's cannon ($\mu_{g,t}$), using Bayesian updating of the local means based on observed samples:

$$P(\mu_{g,t}|x_{t,}x_{1:t-1}) = N\left(\sigma_{L0}^2\left(\frac{\mu_0}{\sigma_{L0}^2}+\frac{\bar{x}}{\sigma^2/t}\right),(\frac{1}{\sigma_{L0}^2}+\frac{1}{\sigma^2/t})^{-1}\right)$$

Where μ_0, σ_{L0}^2 are the prior beliefs on the local mean and variance for each game, and $\bar{x}, \sigma^2$ are the local mean and variance of observed samples 1:t . For the first game, μ_0 is initialized randomly to a position close to the middle of the screen (the same as the initial prior for the spaceship's position, or the global mean), and σ_{L0}^2 is initialized randomly between 5 and 10. For the other games, the model uses estimates from the global mean distribution.

The model tracked its own uncertainty about the estimated local means as the variance of the estimated local mean distribution. With more observed samples (larger *t* in a game), that uncertainty decreased as $\frac{1}{t}\sigma^2$ (see Figure 2A), and the probability of accurate aim on each subsequent trial increased (see Figure 2B). At the end of each game, the model updated its estimate of the spaceship's position in a similar manner, using the last estimate of the game's local mean, and also tracked its uncertainty regarding this global mean estimate.

$$P(M_g|\mu_{g,}\mu_{1:g-1}) = N\left(\sigma_g^2\left(\frac{M_0}{\sigma_{G0}^2}+\frac{\bar{\mu}}{\sigma^2/t}\right),(\frac{1}{\sigma_{G0}^2}+\frac{1}{\sigma^2/t})^{-1}\right)$$

After the first game, the current estimate of the spaceship's position (the global mean) served as prior for the location of the cannon in the next game. The optimal place to aim upon encountering each new cannon (before observing any minions it sends) is this estimate of the spaceship position. Indeed, accurately learning the global distribution across games should lead to a gradual improvement in performance on the first trial of each new game (Fig. 2C) – a key signature of integration of evidence from local to global levels. Conversely, failure to accurately integrate local information from each game to update the estimate of the global mean would lead to no performance improvement on the first trial of each new game (and thus players might aim randomly, or based on the last estimate of the local mean only).

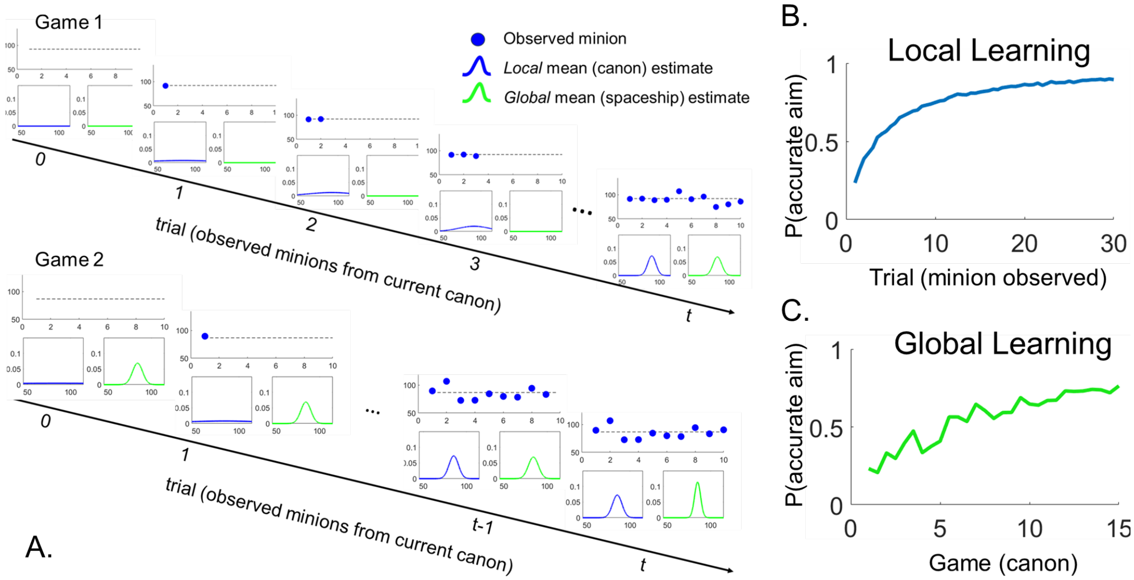

Figure 2: Integrating Local and Global Uncertainty Model. ***A.*** *Example of local observed samples (blue filled circles) changing the estimate of the local mean (blue curves) on each trial of the game; at the end of a game, the final estimated local mean is used as one sample to update the estimate of the global mean (green curve). The local mean estimate resets between games. The global mean estimate carries from game to game, becoming updated at the end of each game.* ***B****. Simulation showing how under the model assumptions, more observed local samples translate to a better estimate of the local mean and improved aim accuracy (i.e., local learning throughout a game).* ***C.*** *Simulated model behavior showing improved aim on the first trial of each new game, i.e. global learning, when the local mean estimates are integrated into estimating the global mean.*

The model assumed that players would aim at the current estimated local mean on each trial (based on the inference processes above), with a noise parameter ε. Under that assumption, it then predicted the decision whether to *stay* or *run*, based on current uncertainty about local and global means. The values for the two actions were computed as follows: *stay* was associated with a reduction in local uncertainty (which decreased as the number of trials in a game increased, leading to diminishing marginal returns on information for spending longer in a game) and an increase in the probability of obtaining local rewards. It was associated with no change in global uncertainty (as the model updated the global estimate only at the end of a game). Conversely, *leave* was associated with an increase in local uncertainty (which reset at the beginning of each game to the current estimate of σ_G^2) and a decrease in estimated local reward, but also with a reduction in global uncertainty, and thus a small increase in the probability of obtaining the larger reward at the end of the task. Thus, the local and global reward estimates traded off on each trial in each game, depending on the uncertainties in the current estimates of the local and global means.

Each action (“Stay” vs. “Leave”) was produced different local and global reward value estimates, which the model used to compute action values.

$$V(stay) \propto (T - t) * (1 - N(0, \sigma_t^2)) + R * (1 - N(0, \sigma_{G-1}^2))$$

$$V(leave) \propto (T - t) * (1 - N(0, \sigma_G^2)) + R * (1 - N(0, \sigma_G^2))$$

It then assumed that participants decided what to do using a softmax rule comparing these two values, with inverse temperature parameter β.

Reinforcement Learning Model

The uncertainty-updating model assumes Bayesian computations of the estimated mean and uncertainty about the mean on each trial (for local means) and at the end of each game (for global mean), as well as forecasting probabilities of reward under different conditions (Stay or Leave) on each trial. While ample evidence exists of the brain engaging in Bayesian updating ([29-31]), including in the learning and decision-making domain ([32; 33]), such a model is expensive in computational and memory resources and requires perhaps more resources than participants are able or willing to expend on the task. A less demanding strategy might entail learning the correct place to aim via trial-and-error, and tracking the change in prediction error (PE, the difference between where the player aimed and where the observed minion appeared on screen) across trials. Previous research strongly indicates that trial-to-trial changes (specifically, the amount of reduction) in prediction error are indicative of the quality of learned information and the need to learn more ([34; 35]). Thus, a model tracking these changes could heuristically estimate, based on thresholds for how much the PE is being reduced, the value of local and global learning, and use that strategy to decide whether to Stay or Leave.

Our reinforcement learning (RL) model updates the local means on each trial based on prediction error PE, with a local learning rate α_L, as

$$\mu_{g,t} = \mu_{g,t-1} + PE * \alpha_L, \text{where}$$

$$PE_t = \ (x_{g,t-1} - \mu_{g,t-1})$$

where $x_{g,t-1}$ is the most recent observed minion position. The global mean M is similarly updated, with a global learning rate α_G. The decision whether to stay or leave a current game is made heuristically, based on the change in tracked local and global prediction errors ΔPE_L, ΔPE_G, and local and global PE-thresholds γ_L, γ_G, as:

$$P(Stay) = \begin{cases} 0, if\ \Delta PE_L < \gamma_L\ and\ \Delta PE_G > \ \gamma_G \\ 1, \qquad\qquad\qquad\qquad\quad otherwise \end{cases}$$

$$\text{with } \Delta PE_L = PE_t - PE_{t-1}$$

As the model learns the underlying distribution of each local game, the local prediction error decreases, and, crucially, the change in prediction error ΔPE_L also decreases; while the local and global PE settle near the level of the generative-process standard deviation (and so never decrease to 0), average ΔPE_L can reach values close to 0 once the model has learned where to aim. A low value of ΔPE_L is indicative that all relevant information from the local game has been learned, and it would be useful to leave and begin a new game. Conversely, a high value for ΔPE_G indicates that there is still relevant information to be learned about the global structure.

Thus, low local PE change and high global PE change indicate that, from an optimal learning perspective, leaving is the correct decision.

Results

As shown in **Table 1**, there were no significant differences between the groups on age, general intellectual ability, gender, or minority status. The OCD group reported a moderate level of OCD symptoms (M = 21.7, SD = 5.1), mild depression (M = 15.3, SD = 1.6), moderate anxiety (M = 16.7, SD = 11.7) as well high levels of intolerance of uncertainty (M = 79.6, SD = 23.3). Fifty-two percent of the OCD participants reporting taking psychiatric medication, including serotonin reuptake inhibitors (n =12), and benzodiazepines (n = 1).

Table 1: Sample demographics

	OCD	HC	*p*
Age	29.4 (9.4)	30.5(10.1)	.702
Female, n (%)	16 (80)	17 (81.0)	.384
Estimated Verbal IQ	120.5 (7.8)	119.2 (6.9)	.567
Employed, n (%)	14 (60.9)	14 (66.7)	.690
Ethnic or racial minority, n (%)	7 (30.4)	7 (33.3)	.837

Local Learning Both the group of patients (OCD) and the group of healthy volunteers (HC) learned the position of the canon in each game, gradually improving their aim throughout the duration of the game. A mixed ANOVA looking at average aim accuracy across Time (20 trials) and Group (OCD and HC) showed a significant main effect of Time, with accuracies later in the game significantly higher than earlier in the game ($F(19, 779) = 54.04, p < 0.01$), and a marginal effect of group, with the average accuracy overall in a game slightly higher for the healthy volunteers than the OCD patients ($M_{HC} = .69, SD_{HC} = .15, M_{OCD} = .655, SD_{HC} = .12, F(1,41) = 4.43, p = 0.042$). There was no significant interaction, suggesting that both groups learn similarly across the trials of a game ($F(19,779)$=1. 043, p = .12; figure 3A).

Integration of Local into Global Learning We examined participants' ability to integrate the information they gained in the local games into their estimate of the global mean by looking at their aim accuracies on the first trial of each new game. As the position of their archer was reset to a random point at the beginning of each new game, and the local means of all games were independently sampled from the global mean, any gain in accuracy on the first trial across games was likely due to participants applying what they learned of the global structure to select a better prior for their estimate of the yet-unseen local structure.

As predicted by our conceptual model of OCD as a disruptor to the local/global integration of information, only the healthy volunteers group showed significant improvements in the first trial across games, while the OCD group had no effect. (Fig. 3B; $M_{HC} = .48, SD_{HC} = .07,\ M_{OCD} = .29,\ SD_{HC} = .072$, significant Group x Time interaction, $F(19,342)=2.01$, $p = 0.007$).

We also examined the accuracy of the participants' guess on the last trial of the task (when they gave an explicit estimate of the global mean by aiming a last arrow at the spaceship). Consistent with reduced global learning, the HC group was slightly more accurate in their estimates ($M = 7.73, SD = 3.50$) than the OCD group ($M = 10.07, SD = 3.95$; two-way independent t-test, $t(42) = 2.06$, $p = .045$). Moreover, the HC group showed a significant correlation between the amount of integration of local into global learning (as measured by the slope of the global learning curves in Fig. 3B) and their accuracy on that final global estimate ($r(19) = 0.419$, $p = 0.01$). This correlation was not present in the OCD group (Figure 3C; $r(19) = 0.02$, $p = 0.59$).

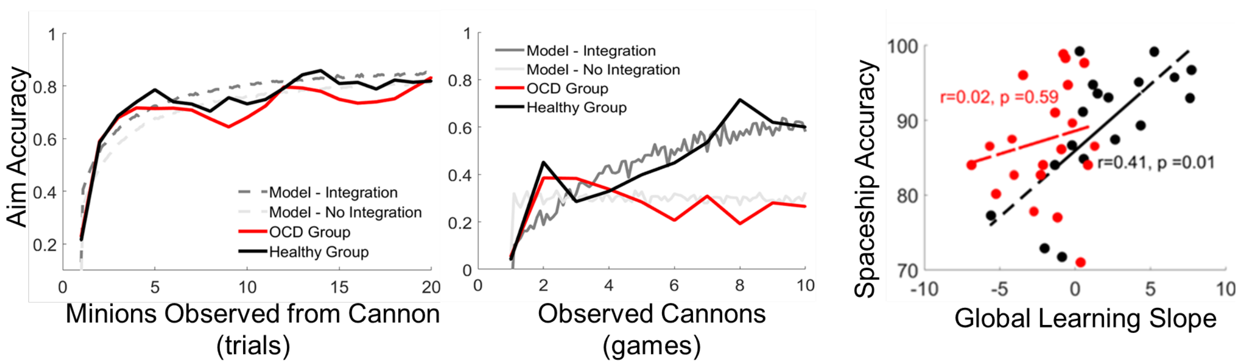

Figure 3: Local Learning and integration of local into global learning. A. Both the control group (black line) and the OCD group (red line) reduced their uncertainty about canon position with more minions observed from one canon. The grey dotted lines show model predictions for this local learning with and without integration of local into global estimates. B. The control group (black line) shows improving aim on the first trial of new games, across the task. The OCD group (red line) does not. The grey lines show the model predictions in the case of integrating information from local to global (dark grey) versus not integrating that information (light grey). C. Global accuracy correlates with the slope of global learning only in the HC group (black markers) and not in the OCD group (red markers).

Local Learning, Global Learning, and Stay/Leave Behavior

We examined trends in participants' stay times (i.e., how many trials they spent in a game before pressing the RUN button to leave) over the course of the task as another behavioral indicator of their ability to integrate local and global learning. Both models predicted shorter stay times in earlier games, when there was higher uncertainty about the global structure, and progressively longer stay times in later games, as more local cannons were observed and thus the uncertainty about the spaceship position was reduced (see Figure 4A, dotted lines).

Participants in the HC group also followed this trend, with average stay times in early games ($M_{game1} = 22.6, SD = 7.4$) significantly different from later games ($M_{game7} = 27.4, SD = 5.82$; $\text{repeated} - \text{measures ANOVA}, \text{F}(6{,}228 = 4.88, \text{p} < 0.01)$). In the OCD group, however, stay times did not show the same pattern, with early and late games both showing significantly longer stay times compared to the HC group ($M_{game1} = 28.2, SD = 4.00, M_{game7} = 29.9, SD = 0.44$, $F(1,228)=2.85$, $p = 0.01$. See Figure 4A, solid lines).

Model fits for learning rates and thresholds from the RL model showed significant differences in parameters for local and global learning, with global learning rates for the OCD group significantly lower than rates for the control group ($M_{Global\_HC} = 0.19, SD = 0.08,$ $M_{Global\_OCD} = 0.07g, SD = 0.07,$ $t(19) = 4.80$, $p < 0.01$; see Figure 4B). Local learning rates, as well as thresholds for local and global learning were not significantly different ($t(19) = 0.85$, $p = 0.4$).

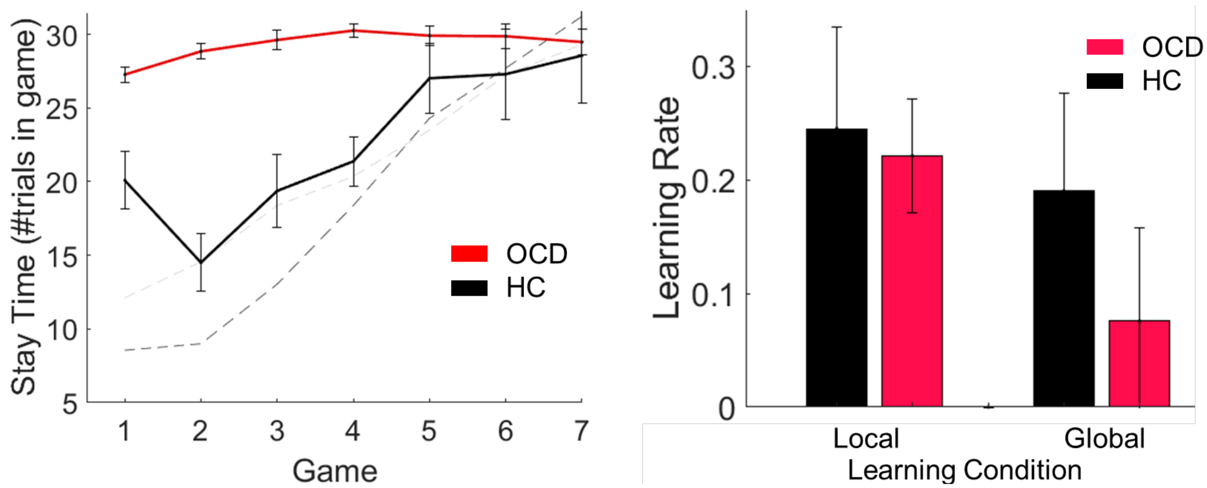

Figure 4: Behavior & Model parameters reflecting global learning. A. Stay times (how many guards attempted per game) increase with number of games in the HC group (black solid line), but not in the OCD group (red solid line). Grey dotted lines represent optimal stay times under different model frameworks (light-grey: RL model; dark grey: uncertainty-updating model). B. Fit model parameters for local and global learning rates in the RL model.

Potential Alternative Strategies for OCD Group's Global Learning To better understand the potential strategies and differences in the local and global learning performance of the OCD group, we tested a series of potential alternatives for their learning, specifically focusing on their strategy for aiming on the first trial of a new game.

We first ran general linear regression model examining the impact of various measures on the first-trial aim position. These measures included: the position of the last observed minion in the previous game (LO), the aim on the last trial of the previous game (LA), the aim on the first trial of the previous game (LFA), the random starting position the task automatically set the archer to on the respective current game (SP), the mean of last three or the last five observed minion positions in the previous game (LM3/LM5), as well as a weighted mean of all observed minion positions so far (with the weight decayed as a function of how many trials back they were observed, AM). Figure 4A shows the results, with none of the regression coefficients significant predicting the first-trial aim.

We also examined the aim-noise parameter ε for differences between groups (Figure 4B). A two-way independent t-test showed a marginal difference ($t(42)=1.95$, $p = 0.057$), with the average aim noise slightly larger in the OCD group ($M_{HC} = 6.92, SD_{HC} = 1.19$; $M_{OCD} = 7.56,$ $SD_{HOCDC} = .098$). The temperature parameter in the softmax did not differ significantly between groups ($t(42) = 0.95$, $p = 0.38$).

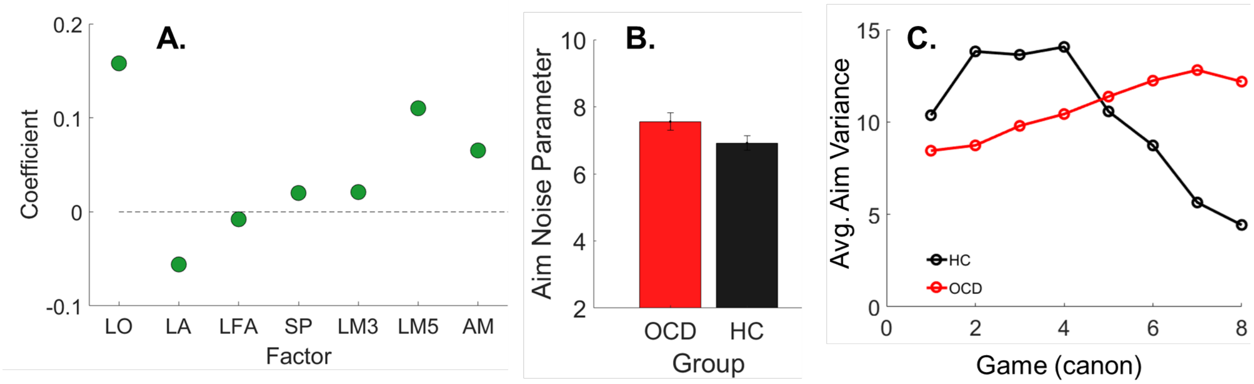

Figure 5: Alternative strategies to integration of local into global learning for the OCD Group. ***A****. No significant coefficients from a general linear regression model when the last (or last few) observed minions (LO, LM3, LM5), the first/last aims on the previous game (LA, LFA), the current archer position (SP), or a weighted mean of all observed minions (AM) were used as potential predictors.* ***B.*** *The model parameter governing the noise in aim is slightly higher for the OCD (red bar) than the HC (black bar) group.* ***C****. The HC group (black) reduces aim variability across games; the OCD group (red) does not.*

To better understand potential differences in aim strategies, we also compared the individual variability in aim across games in the two groups. As participants observe more canons, if they integrate information from local games into their estimate of global structure they should gradually show reduced variability in their aims (as their accuracy improves and they should also be less sensitive to occasional outlier minion positions). Figure 4C shows this pattern of decreasing aim variability in the HC group, but not in the OCD group.

Discussion

Our results suggest the OCD group shows preserved ability to reduce uncertainty at a local, but not global level. This is consistent with existing literature regarding oversampling/overfocusing in OCD [13; 37-38]; however, our task and model further explain conflicting findings on whether OCD patients can accurately represent and reduce uncertainty [16; 17]. Specifically, we propose based on this data that, while the circuitry of reward and uncertainty computations may function to a large degree as it does in a healthy population—as shown by the intact learning at the local level in each game (Fig. 3A), the mechanism responsible for integrating across timescales and generalizing from local to global goals is impaired in OCD.

Further work is needed to precisely characterize the degree and neural bases of this impairment. One possibility involves disrupted reinforcement learning (RL) within the nested hierarchical structure among corticostriatal circuits, where more anterior frontal regions represent global abstract rules that contextualize action selection in posterior circuits involved in achieving subgoals ([39-41]). While RL might be intact at any individual level, impairment in cross-level integration would hamper abstract goals even if the subgoals have been achieved. Another possibility includes an aversion to uncertainty in identifying hidden task states at the level of orbitofrontal cortex ([42]).

Crucially, our task helped identify differences in the ability to process information across local and global levels simultaneously: the OCD group did not show evidence of employing learned global structure to improve local first-trial aims across games (Fig. 3B), nor did they adjust aim variability in later games as the model using information-integration would predict (Fig. 4C). Model simulations show this behavior is consistent with a lack of integration across games of learned local structure in the service of learning the global structure; it is also consistent with the failure to employ the knowledge about the global structure to further guide or refine behavior at the local level.

Interestingly, the performance of the OCD group on the last trial of the task, when participants explicitly estimated the global mean by aiming an arrow at the spaceship, was above chance and only slightly worse than that of the HC group. The lack of improvement on the first trial of each new game (Fig. 3B), as well as the lack of correlation between the spaceship aim and the change in first-trial aim across games (Fig. 3C) suggest that the OCD group did not rely on integration of local information learned throughout the task to estimate the spaceship position. Due to relatively small variances in the generative distributions, it is possible that they could still achieve accuracy on the spaceship trial by treating it as another "local" game. The regression in Figure 4A did not find first or last aims of the previous local game a reliable predictor for aims in the OCD group; however, as there were over 200 trials before this last trial, it is possible players were relying on other parts of their observed local structure that were not tested in the regression model.

While our task design and model allowed testing for local-to-global integration and how information was used simultaneously in the service of multiple goal levels, we could not test, for instance, whether the OCD group integrated local information throughout the task to estimate the global mean but did not apply their knowledge of the global structure to aid performance on local games. This latter would be consistent with findings of impaired transfer of information in OCD [25]. It is also possible that, if this group did to some extent use observed minions to estimate the spaceship's position, they did not perform that computation until the last trial, when instructions explicitly required it (all participants saw a screen warning that they were about to confront the spaceship). There is evidence that OCD patients perform implicit learning tasks better with additional explicit instruction [43; 44]. Our task instructions explicitly mentioned the hierarchical structure (and the benefit of learning the global mean), and all participants reported in post-task debriefings that they were aware of the structure; but there were no reminders through the task of the implicit global learning necessary for optimal performance (as done, for instance, in [43]). It is therefore possible that the OCD group was initially aware of the need for tracking the global and local means simultaneously, but a combination of lower attentional control [45] and the lack of explicit priming during the task hampered their ability to focus on global learning.

It is also possible that the OCD participants did not integrate the local information at all to estimate the global mean due to reduced confidence in their estimates. Ample previous research has found that OCD impairs confidence in one's memory and learning [46-49]; compellingly, this reduced confidence in what one has learned has been found to link to compulsiveness scores, but not general psychiatric status or anxiety scores [50], suggesting that it is indeed linked to the mechanism that drives formation of compulsions. Our current design did not allow participants to observe unlimited minions in any patch and possibly forced them to move on before they were confident that they'd learned an accurate canon position; this may have led them to not fully rely on their observations and inferences to estimate the global structure. One way to test this would be to remove the limitations on the number of arrow and the maximum number of minions from each canon; such alternative design might compensate for the low confidence in one's local estimates enough to improve the integration into global estimates.

Conclusion

Using a hierarchical, predictive inference learning task that involves integrating global goals and local subgoals, we propose that compulsive behaviors associated with OCD are linked to deficits in uncertainty-processing; but, rather than a general failure in uncertainty reduction, OCD impacts the ability to recruit uncertainty pertaining to local subgoals in the service of reducing uncertainty about larger, more general goals. Such a deficit would reconcile existing accounts, by sparing local-level uncertainty reduction (which would show up as intact learning in a variety of tasks) but hampering goal-directed behavior in environments in which subgoals and goals coexist.

References

[1] Guzick, A. G., Candelari, A., Wiese, A. D., Schneider, S. C., Goodman, W. K., & Storch, E. A. (2021). Obsessive–compulsive disorder during the COVID-19 pandemic: a systematic review. *Current psychiatry reports*, *23*, 1-10.

[2] Tanir, Y., Karayagmurlu, A., Kaya, İ., Kaynar, T. B., Türkmen, G., Dambasan, B. N., ... & Coşkun, M. (2020). Exacerbation of obsessive compulsive disorder symptoms in children and adolescents during COVID-19 pandemic. *Psychiatry research*, *293*, 113363.

[3] Silverman, M. E., Nag, S., Kalishman, A., Cox, P. H., & Mitroff, S. R. (2022). Increases in symptoms associated with obsessive-compulsive disorder among university students during the COVID-19 pandemic. *Journal of American College Health*, 1-7.

[4] Boisseau, C. L., Schwartzman, C. M., & Rasmussen, S. A. (2017). Quality of life and psychosocial functioning in OCD. *Obsessive-compulsive disorder: phenomenology, pathophysiology, and treatment. Oxford University Press, New York*.

[5] Nielen, M. M., Den Boer, J. A., & Smid, H. G. O. M. (2009). Patients with obsessive–compulsive disorder are impaired in associative learning based on external feedback. *Psychological Medicine*, *39*(9), 1519-1526. 1

[6] Seow, T. X., Benoit, E., Dempsey, C., Jennings, M., Maxwell, A., O'Connell, R., & Gillan, C. M. (2021). Model-based planning deficits in compulsivity are linked to faulty neural representations of task structure. *Journal of Neuroscience*, *41*(30), 6539-6550.

[7] Zainal, N. H., Camprodon, J. A., Greenberg, J. L., Hurtado, A. M., Curtiss, J. E., Berger-Gutierrez, R. M., ... & Wilhelm, S. (2023). Goal-Directed Learning Deficits in Patients with OCD: A Bayesian Analysis. *Cognitive Therapy and Research*, 1-12.

[8] Gillan, C. M., Morein-Zamir, S., Urcelay, G. P., Sule, A., Voon, V., Apergis-Schoute, A. M., ... & Robbins, T. W. (2014). Enhanced avoidance habits in obsessive-compulsive disorder. *Biological psychiatry*, *75*(8), 631-638.

[9] Soref, A., Dar, R., Argov, G., & Meiran, N. (2008). Obsessive–compulsive tendencies are associated with a focused information processing strategy.*Behaviour research and therapy*, *46*(12), 1295-1299.

[10] Gillan, C. M., Papmeyer, M., Morein-Zamir, S., Sahakian, B. J., Fineberg, N. A., Robbins, T. W., & de Wit, S. (2011). Disruption in the balance between goal-directed behavior and habit learning in obsessive-compulsive disorder. *American Journal of Psychiatry*, *168*(7), 718-726.

[11] Goodwin, G. M. (2022). The overlap between anxiety, depression, and obsessive-compulsive disorder. *Dialogues in clinical neuroscience*.

[12] Rasmussen, S. A., & Tsuang, M. T. (1986). Clinical characteristics and family history in DSM-III obsessive-compulsive disorder. *The American journal of psychiatry*, *143*(3), 317-322.

[13] Hauser, T. U., Moutoussis, M., Nspn Consortium, Dayan, P., & Dolan, R. J. (2017). Increased decision thresholds trigger extended information gathering across the compulsivity spectrum. *Translational psychiatry*, *7*(12), 1296.

[14] Wilson, R. C., Geana, A., White, J. M., Ludvig, E. A., & Cohen, J. D. (2014). Humans use directed and random exploration to solve the explore–exploit dilemma. *Journal of experimental psychology: General*, *143*(6), 2074.

[15] Wilson, R. C., & Niv, Y. (2012). Inferring relevance in a changing world. *Frontiers in human neuroscience*, *5*, 189.

[16] Jacoby, R. J., Szkutak, A., Shin, J., Lerner, J., & Wilhelm, S. (2023). Feeling uncertain despite knowing the risk: Patients with OCD (but not controls) experience known and unknown probabilistic decisions as similarly distressing and uncertain. *Journal of Obsessive-Compulsive and Related Disorders*, *39*, 100842.

[17] Morein-Zamir, S., Shapher, S., Gasull-Camos, J., Fineberg, N. A., & Robbins, T. W. (2020). Avoid jumping to conclusions under uncertainty in Obsessive Compulsive Disorder. *PloS One*, *15*(1), e0225970.

[18] Vaghi, M. M., Moutoussis, M., Váša, F., Kievit, R. A., Hauser, T. U., Vértes, P. E., ... & Dolan, R. J. (2020). Compulsivity is linked to reduced adolescent development of goal-directed control and frontostriatal functional connectivity. *Proceedings of the National Academy of Sciences*, *117*(41), 25911-25922

[19] Gillan, C. M., & Seow, T. X. (2020). Carving out new transdiagnostic dimensions for research in mental health. *Biological Psychiatry: Cognitive Neuroscience and Neuroimaging*, *5*(10), 932-934.

[20] Wise, T., & Dolan, R. J. (2020). Associations between aversive learning processes and transdiagnostic psychiatric symptoms in a general population sample. *Nature communications*, *11*(1), 4179.

[21] Bucarelli, B., & Purdon, C. (2016). Stove checking behaviour in people with OCD vs. anxious controls. *Journal of behavior therapy and experimental psychiatry*, *53*, 17-24.

[22] Solway, A., Lin, Z., & Vinaik, E. (2021). Transfer of information across repeated decisions in general and in obsessive–compulsive disorder. *Proceedings of the National Academy of Sciences*, *118*(1), e2014271118.

[23] Dar, Reuven. "Elucidating the mechanism of uncertainty and doubt in obsessive-compulsive checkers." *Journal of behavior therapy and experimental psychiatry* 35, no. 2 (2004): 153-163.

[24] Dar, R., Sarna, N., Yardeni, G., & Lazarov, A. (2022). Are people with obsessive-compulsive disorder under-confident in their memory and perception? A review and meta-analysis. *Psychological Medicine*, 1-9.

[25] Fradkin, I., Ludwig, C., Eldar, E., & Huppert, J. D. (2020). Doubting what you already know: Uncertainty regarding state transitions is associated with obsessive compulsive symptoms. *PLoS computational biology*, *16*(2), e1007634.

[26] Nassar, M. R., Rumsey, K. M., Wilson, R. C., Parikh, K., Heasly, B., & Gold, J. I. (2012). Rational regulation of learning dynamics by pupil-linked arousal systems. Nature neuroscience, 15(7), 1040-1046.

[27] Geana, A., Wilson, R. C., Daw, N., & Cohen, J. D. (2016). Information-seeking, learning and the marginal value theorem: a normative approach to adaptive exploration. In *Proceedings of the Annual Meeting of the Cognitive Science Society* (Vol. 38).

[28] Huys, Q. J., Maia, T. V., & Frank, M. J. (2016). Computational psychiatry as a bridge from neuroscience to clinical applications. *Nature neuroscience*, *19*(3), 404-413.

[29] Wise, T., Michely, J., Dayan, P., & Dolan, R. J. (2019). A computational account of threat-related attentional bias. *PLoS computational biology*, *15*(10), e1007341.

[30] Yon, D., & Frith, C. D. (2021). Precision and the Bayesian brain. *Current Biology*, *31*(17), R1026-R1032.

[31] Loeb, G. E., & Fishel, J. A. (2014). Bayesian action & perception: representing the world in the brain. *Frontiers in neuroscience*, *8*, 341.

[32] Poudel, G. R., Bhattarai, A., Dickinson, D. L., & Drummond, S. P. (2017). Neural correlates of decision-making during a Bayesian choice task. *NeuroReport*, *28*(4), 193-199.

[33] Fleming, S. M., & Daw, N. D. (2017). Self-evaluation of decision-making: A general Bayesian framework for metacognitive computation. *Psychological review*, *124*(1), 91.

[34] Frömer, R., Nassar, M. R., Bruckner, R., Stürmer, B., Sommer, W., & Yeung, N. (2021). Response-based outcome predictions and confidence regulate feedback processing and learning. *elife*, *10*, e62825.

[35] Sewell, D. K., Jach, H. K., Boag, R. J., & Van Heer, C. A. (2019). Combining error-driven models of associative learning with evidence accumulation models of decision-making. *Psychonomic Bulletin & Review*, *26*, 868-893.

[36] Michely, J., Martin, I. M., Dolan, R. J., & Hauser, T. U. (2023). Boosting serotonin increases information gathering by reducing subjective cognitive costs. *Journal of Neuroscience*, *43*(32), 5848-5855.

[37] Harkin, B., Miellet, S., & Kessler, K. (2012). What checkers actually check: an eye tracking study of inhibitory control and working memory.

[38] Hermans, D., Martens, K., De Cort, K., Pieters, G., & Eelen, P. (2003). Reality monitoring and metacognitive beliefs related to cognitive confidence in obsessive–compulsive disorder. *Behaviour Research and Therapy*, *41*(4), 383-401.

[39] Frank, M. J., & Badre, D. (2012). Mechanisms of hierarchical reinforcement learning in corticostriatal circuits 1: computational analysis. *Cerebral cortex*, *22*(3), 509-526.

[40] Collins, A. G. E., & Frank, M. J. (2016). Motor demands constrain cognitive rule structures. *PLoS computational biology*, *12*(3), e1004785.

[41] Balleine, B. W., Dezfouli, A., Ito, M., & Doya, K. (2015). Hierarchical control of goal-directed action in the cortical–basal ganglia network. *Current opinion in behavioral sciences*, *5*, 1-7.

[42] Schuck, N. W., Cai, M. B., Wilson, R. C., & Niv, Y. (2016). Human orbitofrontal cortex represents a cognitive map of state space. *Neuron*, *91*(6), 1402-1412.

[43] Soref, A., Liberman, N., Abramovitch, A., & Dar, R. (2018). Explicit instructions facilitate performance of OCD participants but impair performance of non-OCD participants on a serial reaction time task. *Journal of Anxiety Disorders*, *55*, 56-62.

[45] Basel, D., Hallel, H., Dar, R., & Lazarov, A. (2023). Attention allocation in OCD: A systematic review and meta-analysis of eye-tracking-based research. *Journal of Affective Disorders*.

[46] Vaghi, M. M., Luyckx, F., Sule, A., Fineberg, N. A., Robbins, T. W., & De Martino, B. (2017). Compulsivity reveals a novel dissociation between action and confidence. *Neuron*, *96*(2), 348-354.

[47] Wheaton, M. G., Messner, G. R., & Marks, J. B. (2021). Intolerance of uncertainty as a factor linking obsessive-compulsive symptoms, health anxiety and concerns about the spread of the novel coronavirus (COVID-19) in the United States. *Journal of obsessive-compulsive and related disorders*, *28*, 100605.

[48] Morriss, J., McSorley, E., & Van Reekum, C. M. (2018). I don't know where to look: the impact of intolerance of uncertainty on saccades towards non-predictive emotional face distractors. *Cognition and emotion*, *32*(5), 953-62.

[49] Gagné, J. P., & Radomsky, A. S. (2020). Beliefs about losing control, obsessions, and caution: An experimental investigation. *Behaviour research and therapy*, *126*, 103574.

[50] Hermans, D., Engelen, U., Grouwels, L., Joos, E., Lemmens, J., & Pieters, G. (2008). Cognitive confidence in obsessive-compulsive disorder: Distrusting perception, attention and memory. *Behaviour research and therapy*, *46*(1), 98-113.